



Binary Search Trees

Let's say that I have a collection of numbers...




There are a few fundamental
operations I can perform



I can **lookup** an item to see if it exists

lookup(




, 12)



true

I can **lookup** an item to see if it exists

lookup(




, 10)




false

I can **add** an item to the collection





I can **add** an item to the collection


add(



, 42)



For now, I haven't talked about how we actually **implement** the collection




Only it's **specification**


Ask yourself: what's the
simplest possible implementation?


Ask yourself: what's the
simplest possible implementation?

For me that would be a **list**




Implemented as a list...





Note: collection is
unordered!

Implemented as a list...



Can implement these two operations

```
def insert(l,n):  
    return l.insert(0,n) # Extend front
```

```
def lookup(l,n):  
    for i in l:  
        if (i == n): return true # Found!  
    return false
```


Can implement these two operations

$O(1)$

```
def insert(l,n):  
    return l.insert(0,n) # Extend front
```

$O(n)$

```
def lookup(l,n):  
    for i in l:  
        if (i == n): return true # Found!  
    return false
```



Implementation: list

Operations

insert $O(1)$


lookup $O(n)$

Can we do better for lookup?

Answer: **yes**, using slightly smarter data-structure

But all data structures have trade-offs

First some intuition...



When I lookup, I go through list one-by-one...


12

23

42


14

-2




When I lookup, I go through list one-by-one...





The *crucial trick* is to **sort**



Is 12 in this sequence?

Start looking here

-2

12

14

23

42

Is 12 in this sequence?

Start looking here

-2

12

14

23

42

14 > 12, so **I know it must be in lower half**

Is 12 in this sequence?

Keep looking here...

-2

12



14

23

42

To lookup i in a sequence l :


- Start at the middle element, call it j
- If $i = j$, then return true
- If $i < j$, then repeat with the lower half of the list
- If $i > j$, then repeat with the upper half

Logarithmic work! Every step eliminates **half!**

To lookup i in a list l :

- Start at the middle element, call it j
- If $i = j$, then return true
- If $i < j$, then repeat with the lower half of the list
- If $i > j$, then repeat with the upper half

Logarithmic work! Every step eliminates **half!***



Logarithmic work! Every step eliminates **half!***

*But if we use a list, we **don't** get $\log(n)$ lookup!



Lists aren't **random access**, so first step takes linear time!

To lookup i in a sequence:

- Start at the middle element, call it j
- If $i = j$, then return true
- If $i < j$, then repeat with the lower half of the list
- If $i > j$, then repeat with the upper half

Therefore, actually $n!$:(



One possible solution: represent as an **array**



Arrays **are** random access


Lookup now $\log(n)$!

One possible solution: represent as an **array**





Downside: insertions must copy the **whole array**

Insert 13





One possible solution: represent as an **array**





Downside: insertions must copy the **whole array**

Insert 13




Newly allocated space

One possible solution: represent as an **array**



Downside: insertions must copy the **whole array**

Insert 13




Lesson: trade-offs *nuanced*, no silver bullet

Let's apply this intuition to **trees**


BST: Binary Tree that has the...

Binary Search Property


Every item in left child $<$ parent, vice versa




**Everything over here had better be $<$ 14
(Even in children of this node)**




Let's say we want to look for 12...




Start at the top of the tree




$12 < 14$, so look down left child





At a high level: Walk down tree left / right until you find what you're looking for


The **structure** of the tree tells you *which way* to go



Implementing Lookup


```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```

```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```





```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```

Looking for 12




```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```

Looking for 23




```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```


Looking for 13



Now let's say I want to insert 13...




Start at top




$13 < 14$, so must add to left


Start at top




13 > 12, so must add to right



Oops! Nothing here. Add **new** node



Oops! Nothing here. Add **new** node



```
# Assume node(elem,left,right) is a constructor
def add(t,i):
    if t == null: new node(i,null,null)
    if t.elem == i: return t
    else if t.elem < i:
        return node(t.elem,add(t.left,i),t.right)
    else if t.elem > i:
        return node(t.elem,t.right,add(t.right,i))
```

Observation: BSTs can store **more than just numbers**

Only need **total ordering** (any two can be compared)

- ➡ Strings

- ➡ Doubles

- ➡ Other user defined types

- ➡ Some langs allow overloading <

Observation: BSTs can store **more than just numbers**

Only need **total ordering** (any two can be compared)

- ➡ Strings

- ➡ Doubles

- ➡ Other user defined types

- ➡ Some langs allow overloading <

Can also use as basis for other data structures
(e.g., dictionary: nodes key/value pairs)

insert $O(\text{height})$

lookup $O(\text{height})$

$O(\log(\text{size}))$ when *balanced*


Naive insertion **does not** balance tree :(

insert $O(\text{height})$


lookup $O(\text{height})$

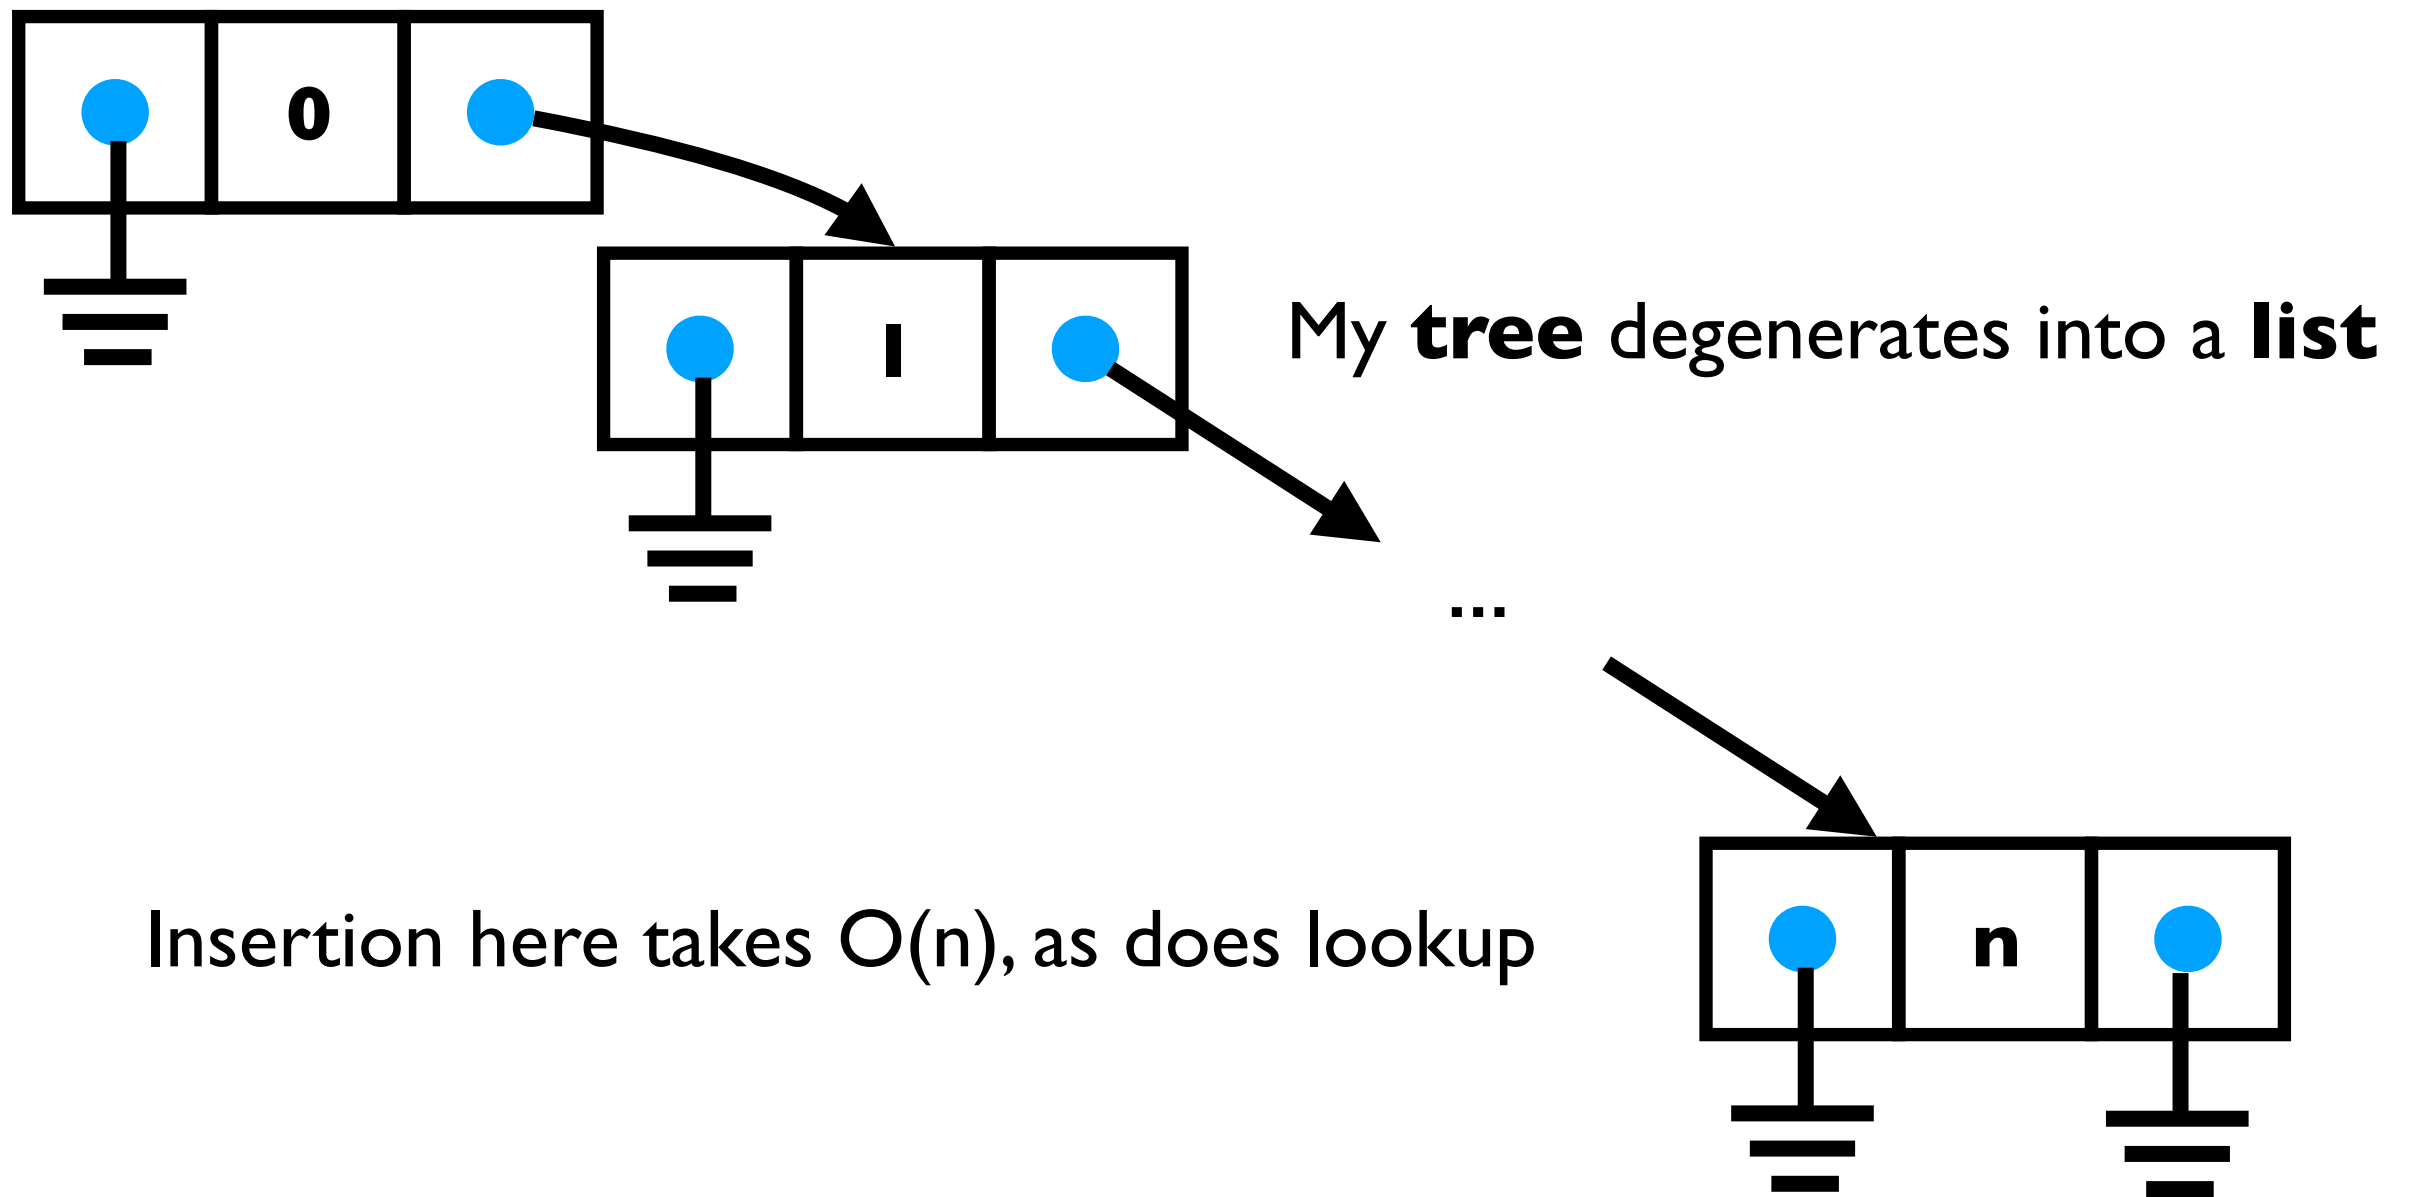
$O(\log(\text{size}))$ when *balanced*

Let's say I start with a 1-element tree...



Then extend it...





Question

Can we ensure good performance generally?

- Precompute **best** BST (dynamic programming)
- Randomize insertion order
- Build *even smarter* data structures:
 - Red-Black trees maintain “balanced-ish” trees
 - AVL trees “rebalance” the tree

(Beyond scope, today)

List

insert

$O(1)$



lookup

$O(n)$



Simple

Insertions frequent

Lookups frequent

List

insert

$O(1)$



lookup

$O(n)$



Simple

Insertions frequent

Lookups frequent

Sorted Array

Also allocates lots of memory

insert

$O(n)$



lookup

$O(\log(n))$



Lookups frequent

Insertions frequent

List

insert

$O(1)$



lookup

$O(n)$



Simple

Insertions frequent

Lookups frequent

Sorted Array

Also allocates lots of memory

insert

$O(n)$



lookup

$O(\log(n))$



Lookups frequent

Insertions frequent

BST

balanced

insert

$\sim O(\log(n))$



Lookups frequent

Insertions frequent

lookup

$\sim O(\log(n))$



Maintaining balance hard