

The Hash Array-Mapped Trie (HAMT)

Kris Micinski

Logistics

- I'm (probably) gone next Tuesday
- This Thursday: course / exam review
- Next Thursday (probably): present projects / competition
- None of the HAMT particulars will be on the exam

HAMT — High-Level Benefit

- **Persistent** hash map / set with:
 - Constant time insertion
 - Constant time lookup
 - Robust cache performance

Problems with other DS


- Lookup / insertion for **balancing** binary trees:
 - $\log(n)$ with **imperative** version
 - Not persistent data structure
- Same thing with hash tables...
- Coming up with persistent hash map is hard!

Motivation: phone books over time...

Keys		Values
John	→	(111) 111-1111
Kelly	→	(222) 222-2222
Sam	→	(333) 333-3333
José	→	(444) 444-4444


Keys

Values



Keys

Values



... Aug 10 Aug 11 Aug 12 ...

Keys	Values
John	→ (111) 111-1111
Kelly	→ (222) 222-2222
Sam	→ (333) 333-3333
José	→ (444) 444-4444

Keys	Values
John	→ (111) 111-1111
Kelly	→ (222) 222-2222
Sam	→ (333) 333-3333
José	→ (444) 444-4444 → (555) 555-5555

Keys	Values
John	→ (111) 111-1111
Kelly	→ (222) 222-2222
Sam	→ (333) 333-3333 → (666) 666-6666
José	→ (444) 444-4444 → (555) 555-5555

Keys

Values

John
Kelly
Sam
José



(111) 111-1111



(222) 222-2222



(333) 333-3333



(444) 444-4444

Keys

Values

John
Kelly
Sam
José



(111) 111-1111



(222) 222-2222



(333) 333-3333




~~(444) 444-4444~~

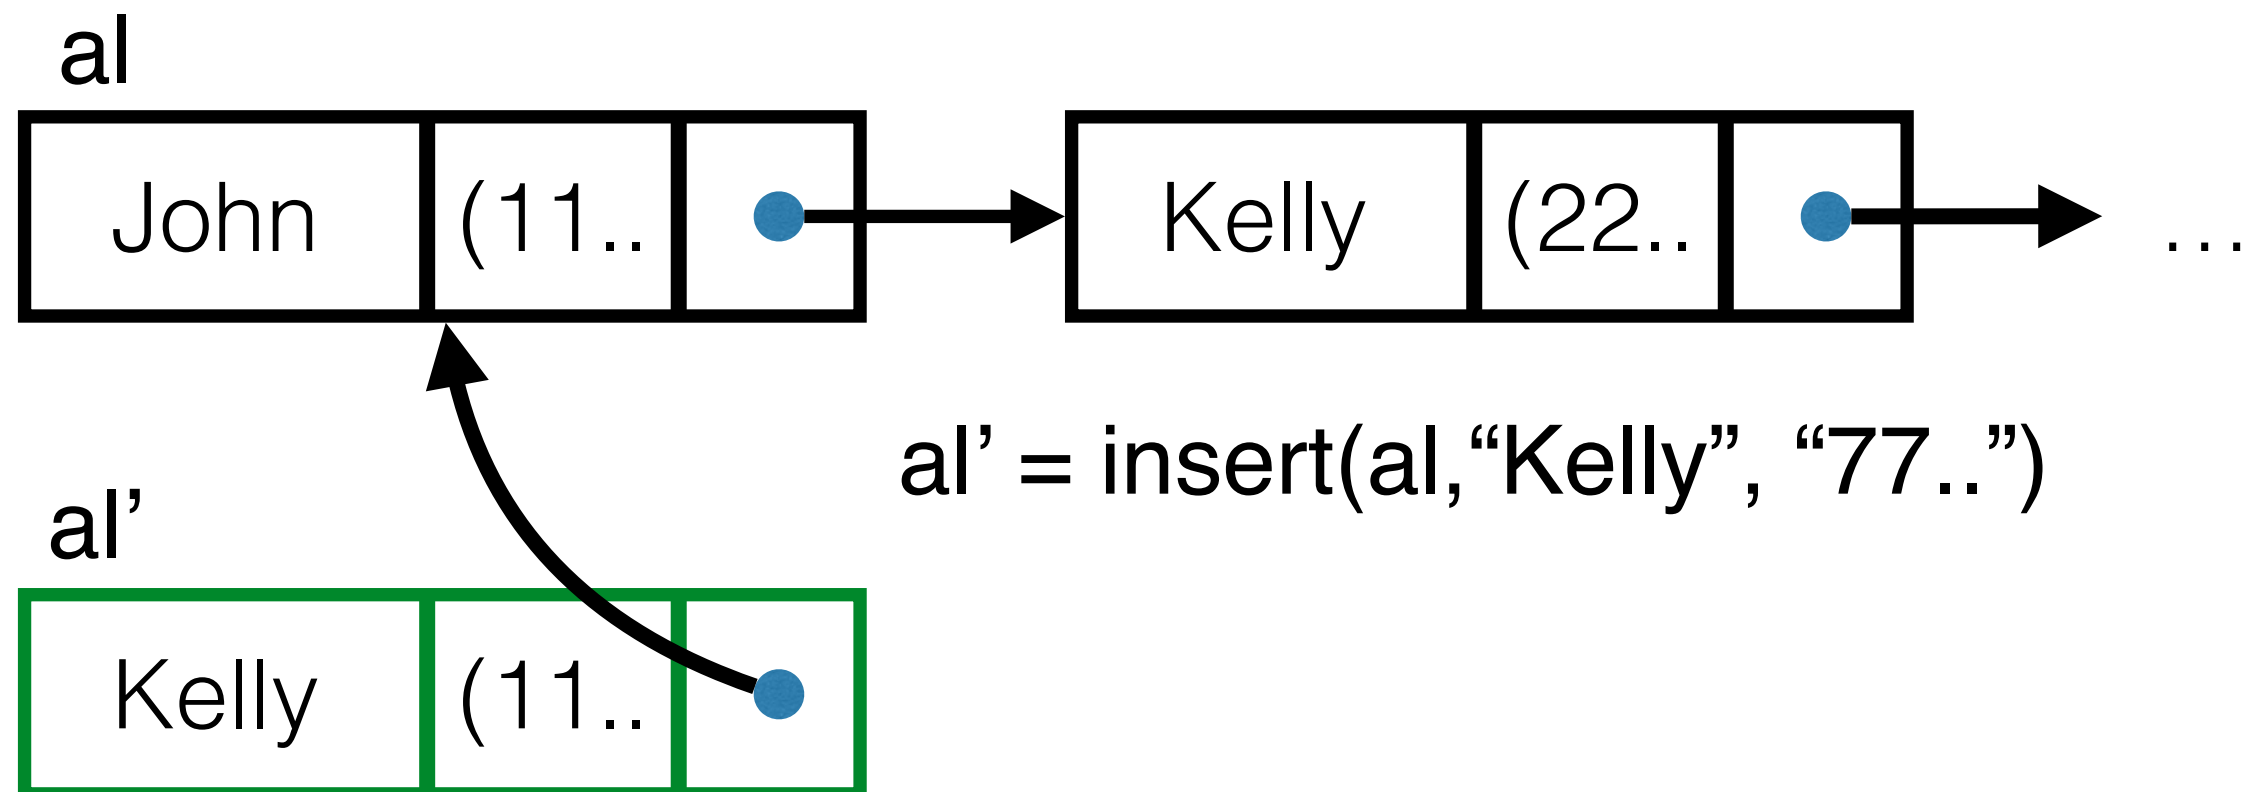
(555) 555-5555

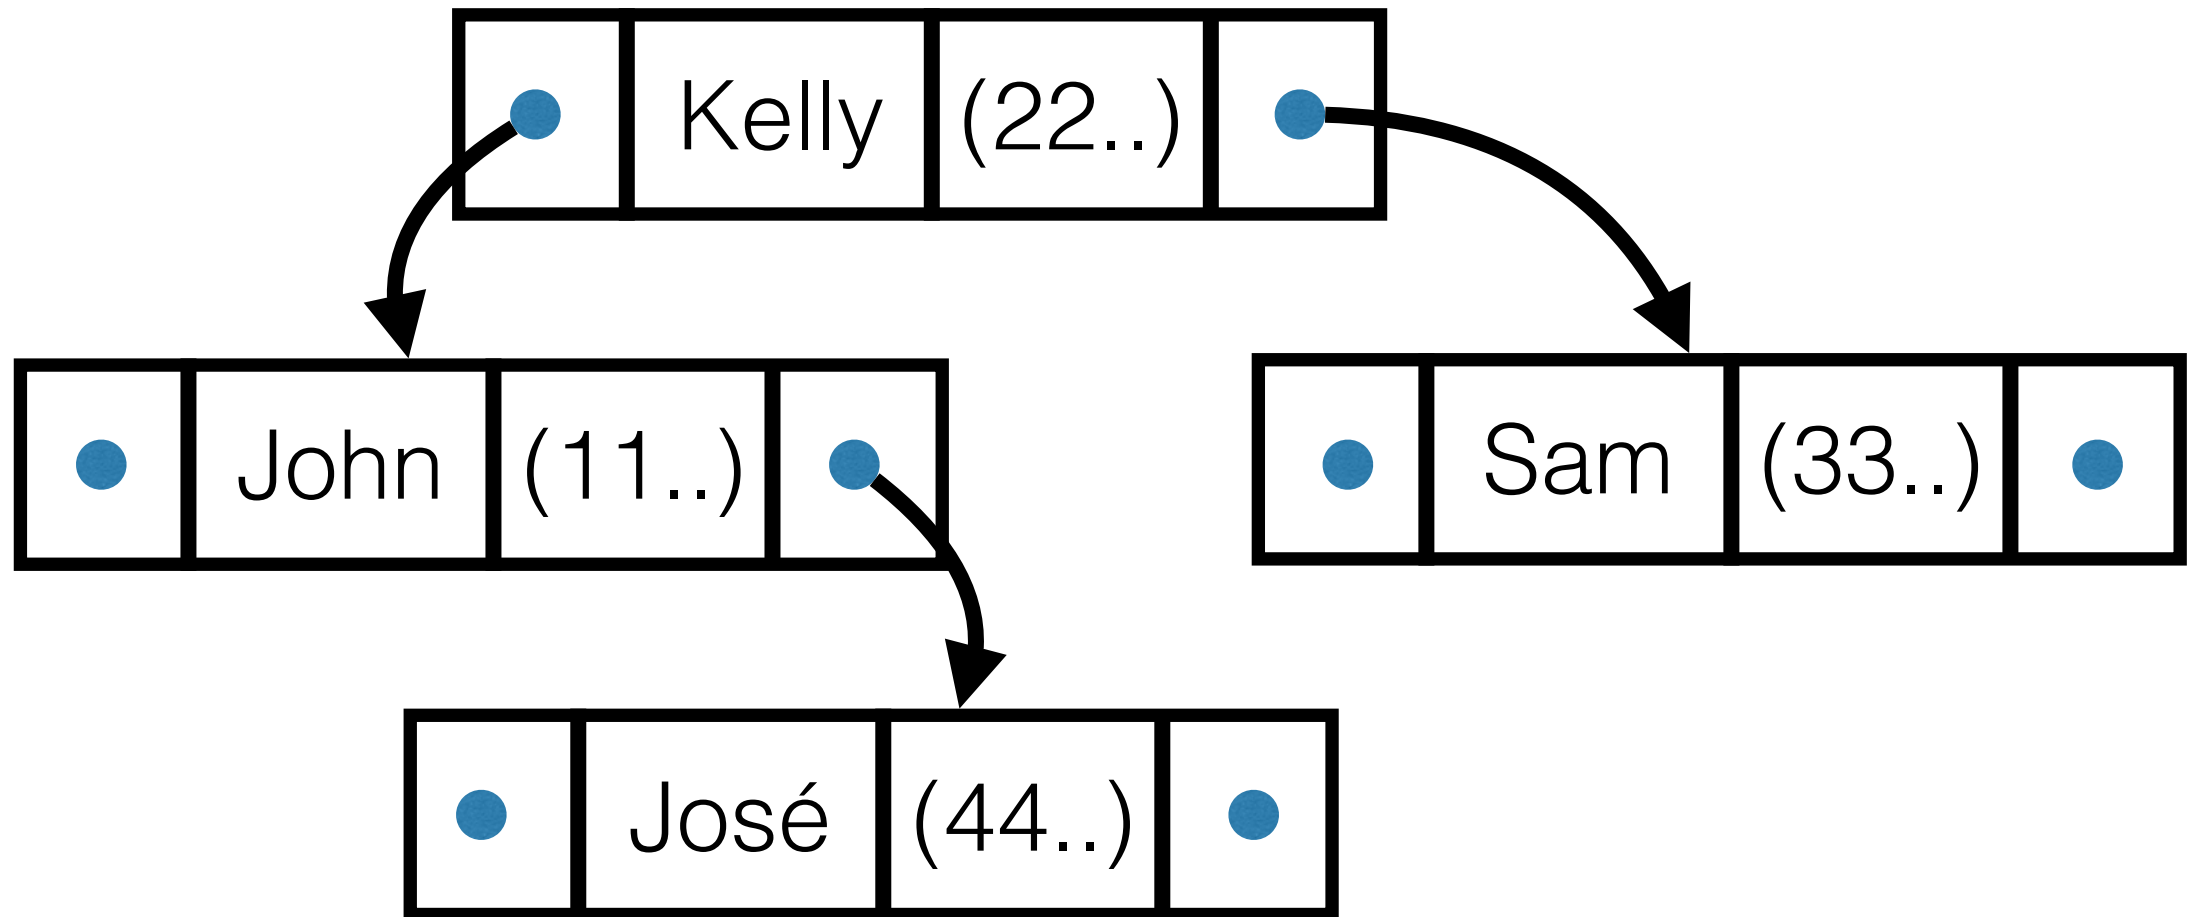
This takes $O(n)$!

copy



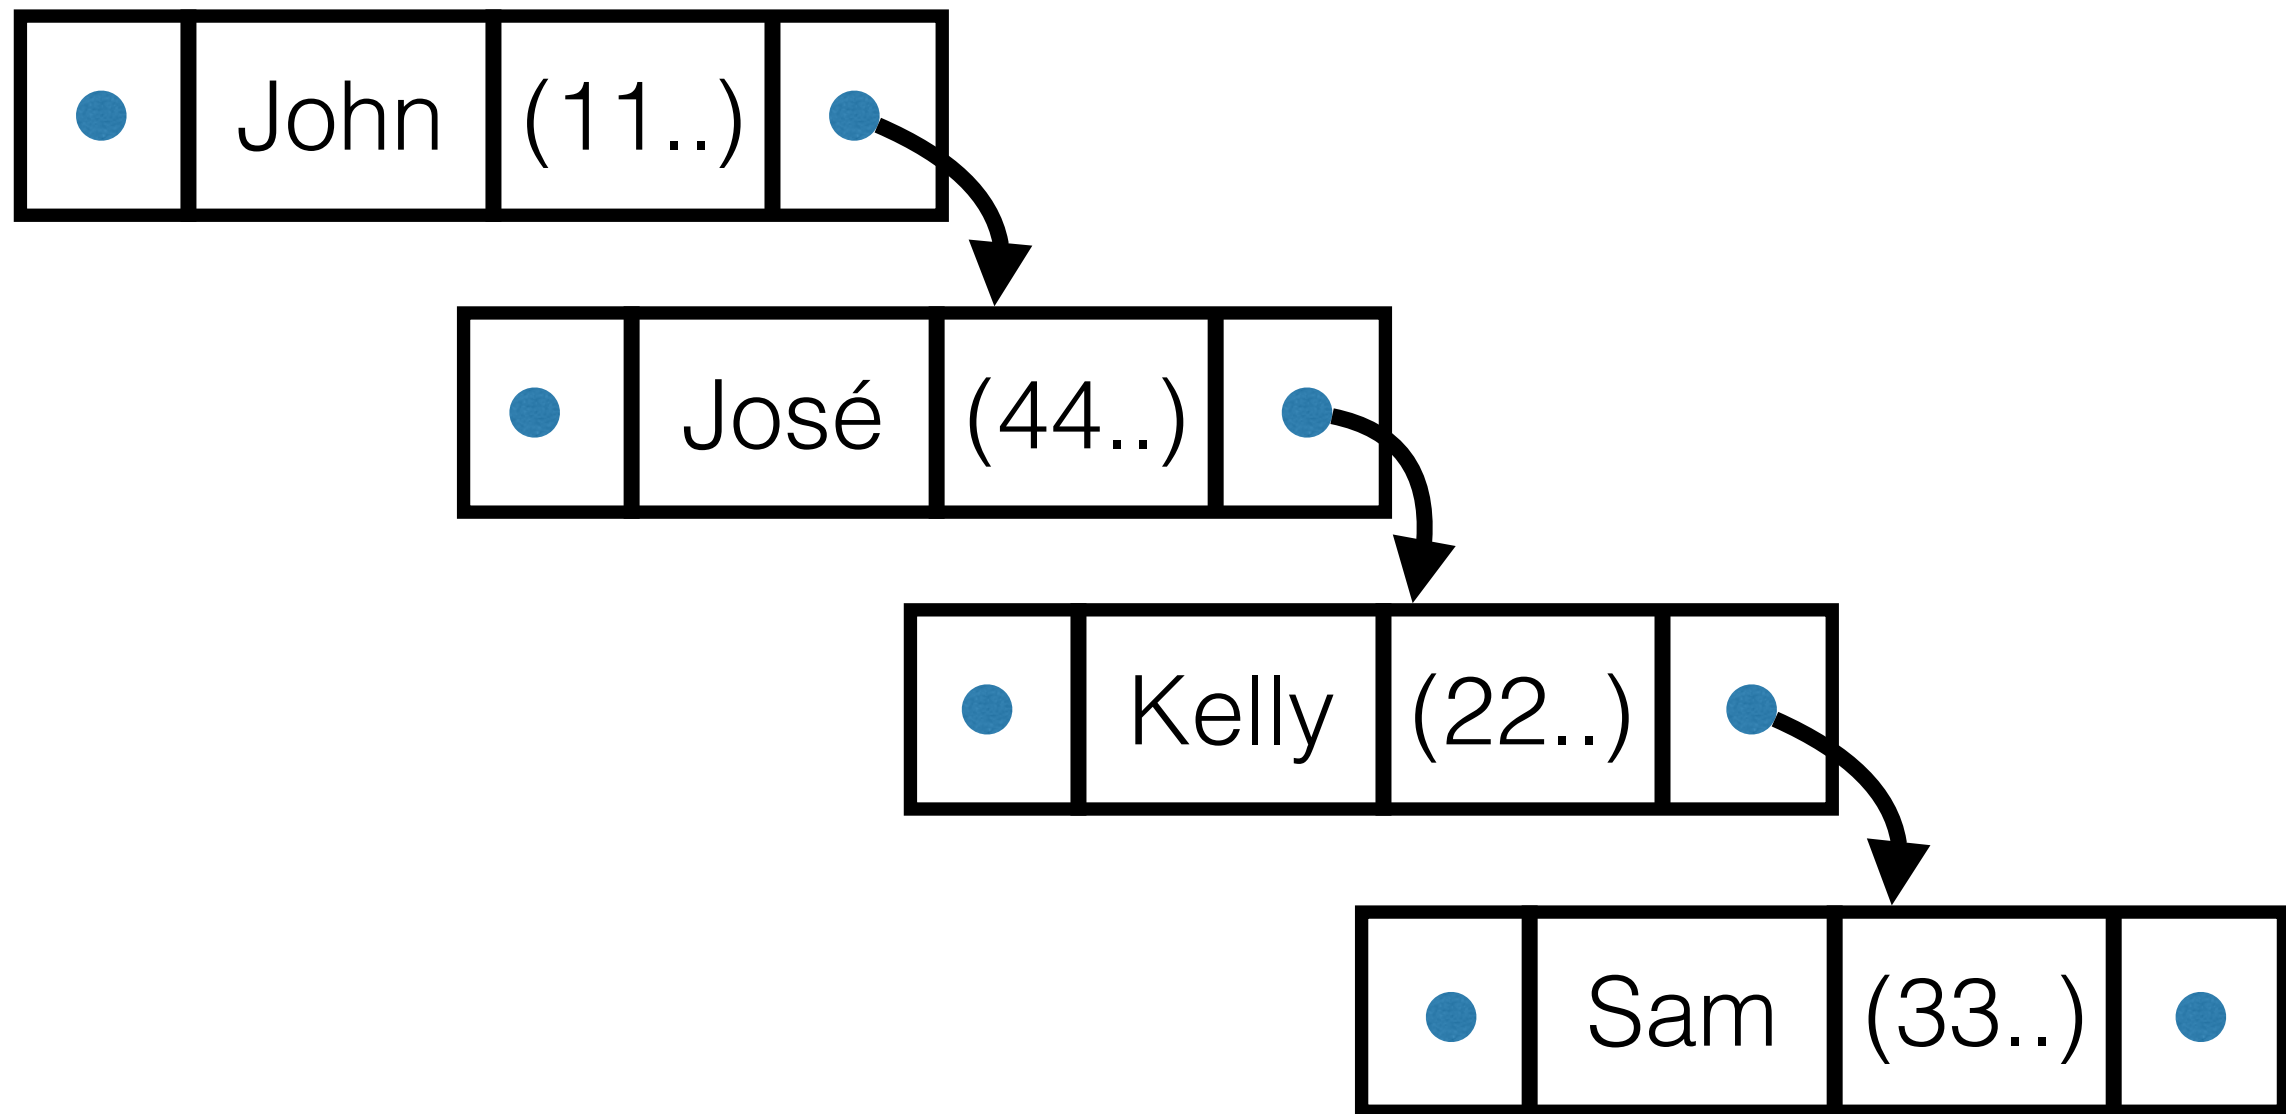
Inserting into association list: either $O(1)$ or $O(n)$
Depends if you want time or space...





Balanced binary trees are good....

But unbalanced binary trees are not...!



(Naive insertion **potentially** leads to $O(n)$!)

What's one way to **approximate**
balanced binary trees?

(I.e., if I want to insert n names and end
up with an “almost” balanced tree?)

Insight: **randomize** insertion order

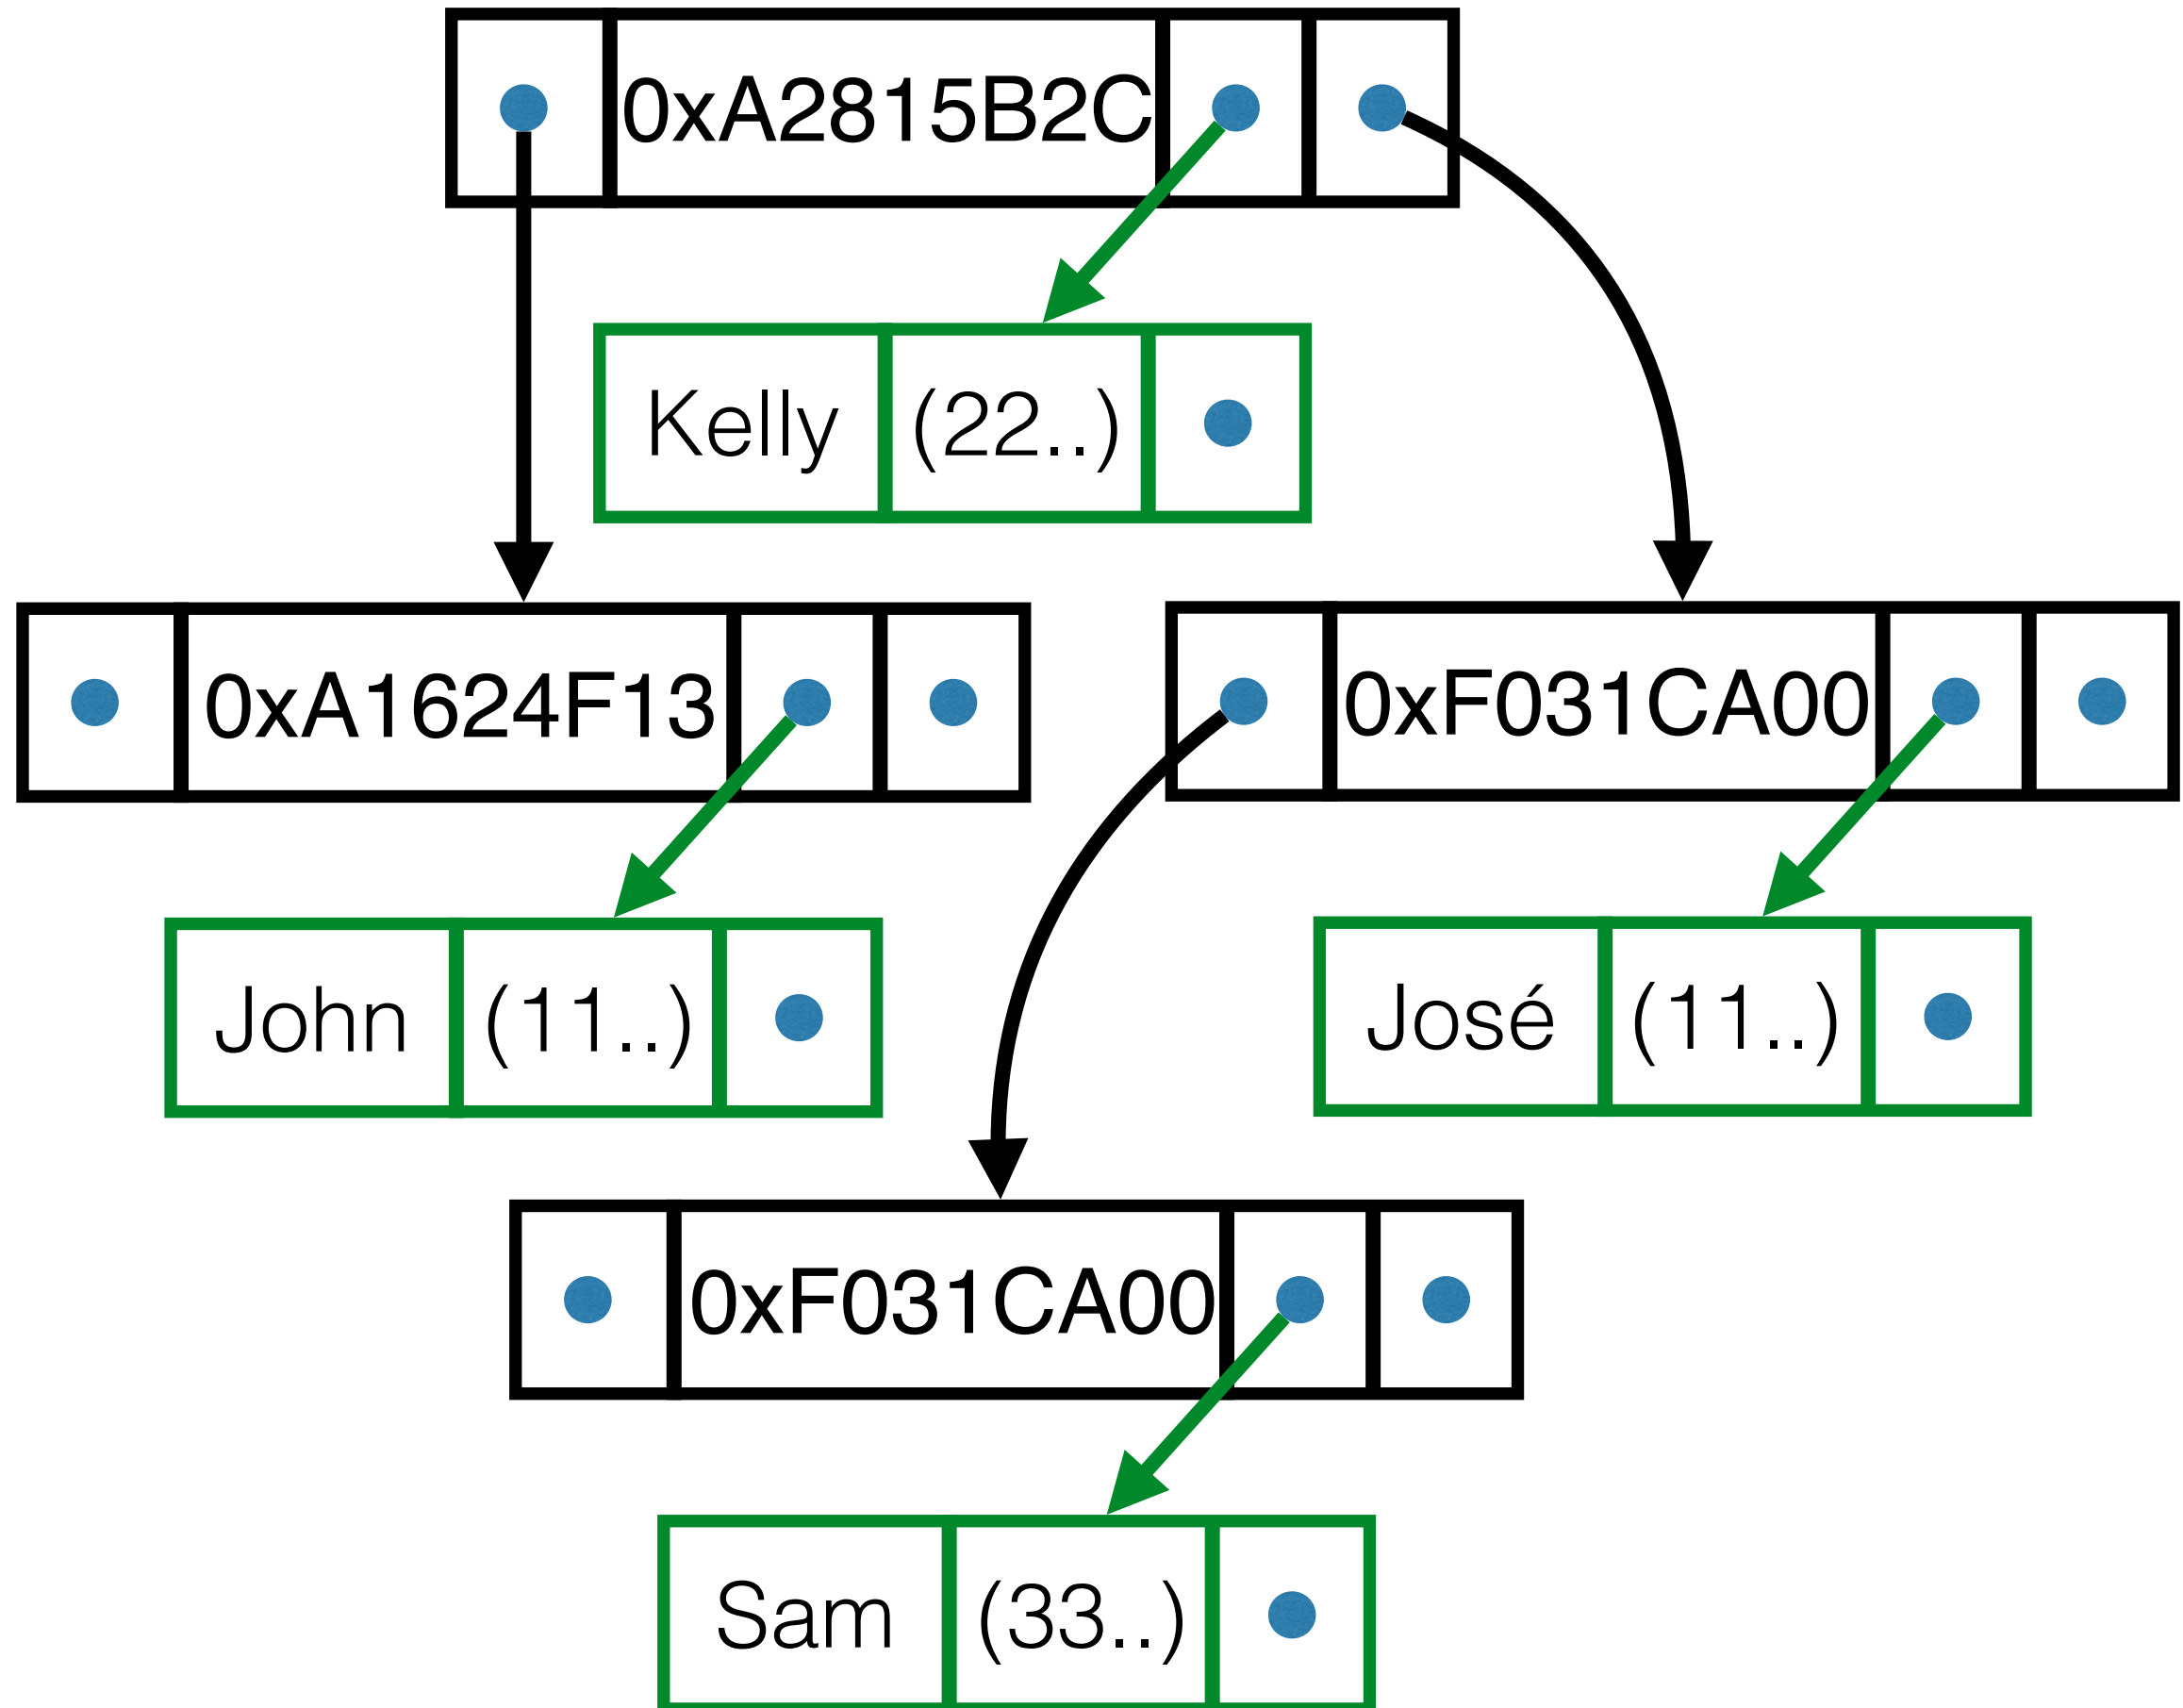
I can do this by storing a **hash**!

Now: each node of our binary tree
stores a **hash** and an **association list**

(Why do you need this?)

Observation 1

Store hashes as keys into a tree, back
it with association list



Observation 2

Instead of tree, why not use a **trie**?

0x100000001



0x100000002



0x100000003




0x100000004




0x100000005



...





**Now each successive bit in the trie indexes
a successive bit in the hash**

(Of course, each leaf still needs to be backed
by assn. list!)

Observation 2

Now lookup takes **at most** 64!
(Because that's how long our hash is!)

Ergo: Insertion takes $\sim O(1)$ time!

Observation 3

64 is still pretty crummy constant factors!


And it **always** takes 64 if we use a trie!

How can we do better..?

How can we do better..?


Idea: store more than 1 bit per node!

Store 64 subtries
per node



(Can vary this up / down)

Now our trie looks like this!



To insert, walk over **chunks** of the hash!

(Walk through on board...)

Observation 3

Get better constant factors by storing more
subtries per node

(Downside: each node takes more memory)

Observation 4

Don't store the **whole** trie until we really **need** to

No other keys could **possibly** be stored in this bucket!

John → 0x00...


Kelly → 0x4A...

Sam → 0xA1...

José → 0xFC...

So don't store a subtrie!

Takes up less memory, also faster!



Question now: how do we implement insert?

Observation: form **new** subtrees in case of collision at some node


insert(map,"Sam", "(33..")

Sam → 0xFFF...

José → 0xFC0...

0xFFF = 111111 111111

0xFC0 = 111111 000000



Split!

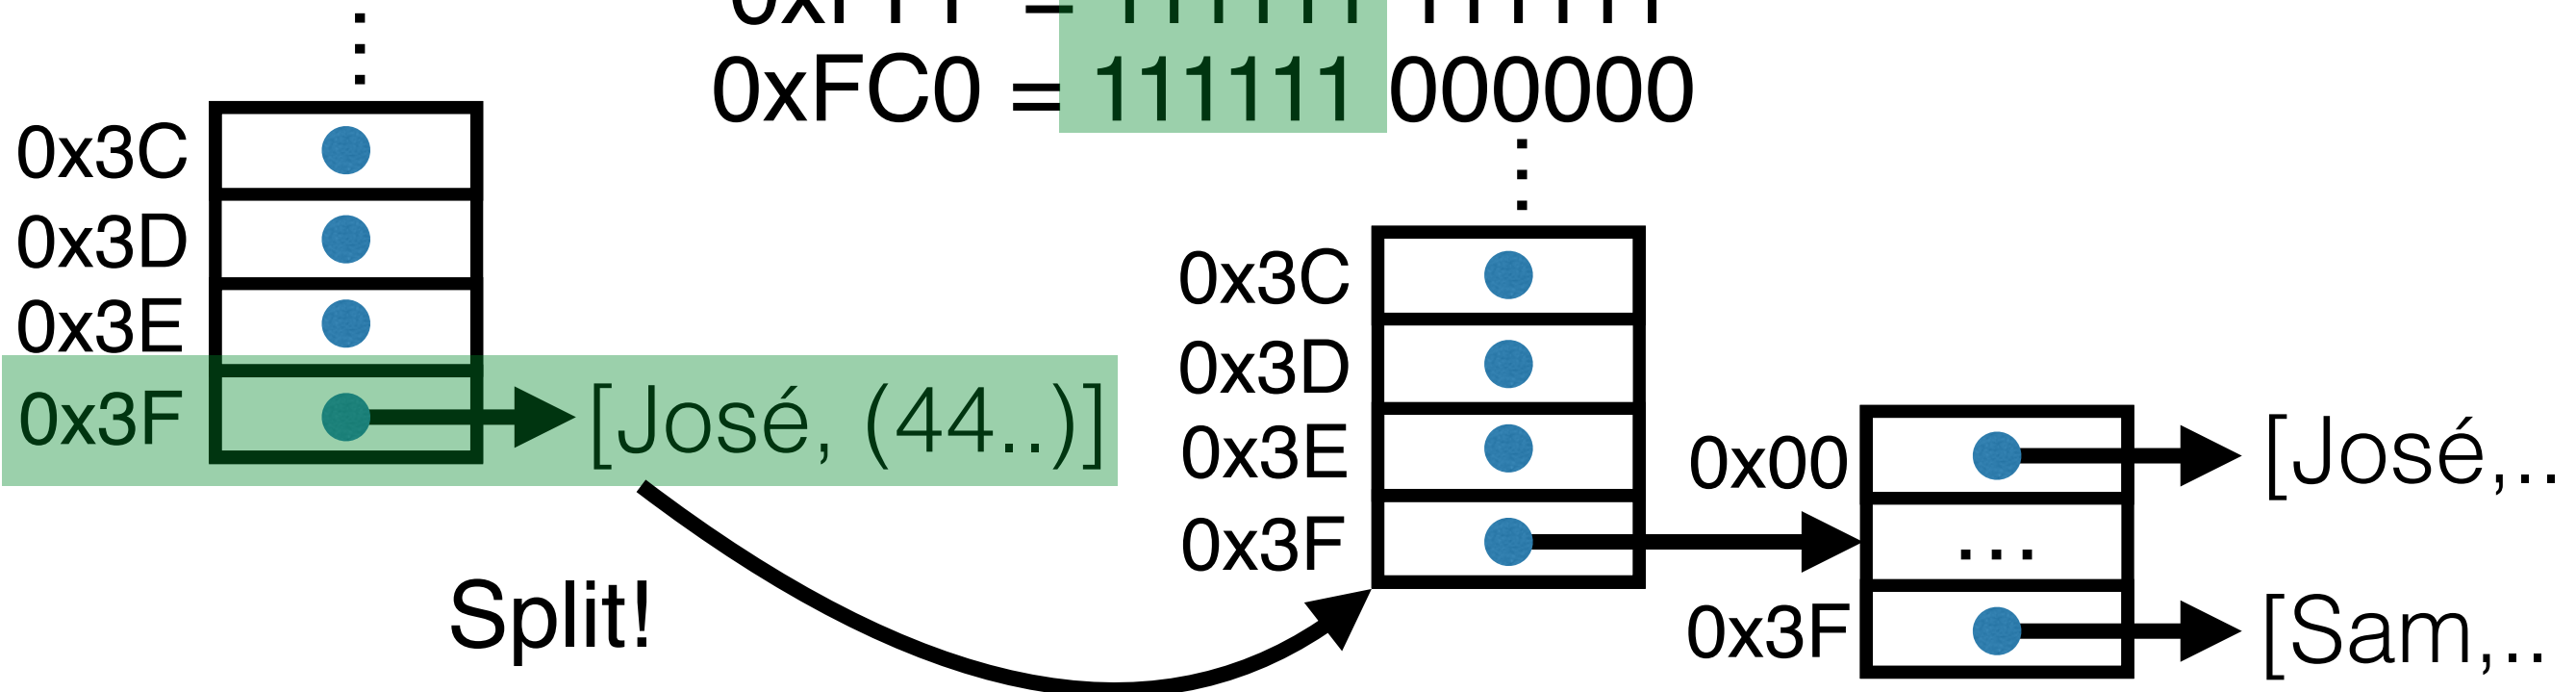
insert(map,"Sam", "(33..")

Sam \longrightarrow 0xFFF...

José \longrightarrow 0xFC0...

0xFFF = 111111 111111

0xFC0 = 111111 000000



Observation 4

Don't store the **whole** trie until we really **need** to

(Reduces space the trie takes up! Storing fewer
keys = fewer places taken up!)

Observation 5


Storing 64 subtries is **still** very expensive!

Solution: store array of subtries, length of array is n where n is number of occupied buckets!


(Kris illustrates on board...)

Idea: use **bitmap**!

Old representation



New representation



Putting it all together

Store hashes as keys into a tree, back it with association list

Instead of tree, use trie

Lower constant factors: store n (= 64 in our ex) subtries
(Con: more space for each subtrie!)

Two clever hacks:

- Don't store subtries until you **absolutely need** to
- Use bitmap to reduce the need for n subtries until needed

Result

- Fast persistent hash map!
- Great constant factors
- Low overhead when it's not storing much stuff
- Higher overhead as it fills up, but never far past constant factor!
- Useful in implementing interpreters, any other place when you need efficient hash map
- Most of the time a regular HT is probably fine, but think of HAMT!