

Implementing




Binary Search Trees & Dictionaries

BST: Binary Tree that has the...

Binary Search Property

Every item in left child $<$ parent, vice versa



**Everything over here had better be $<$ 14
(Even in children of this node)**

Implementing Lookup

```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```

```
# Assume t is a tree with .left, .right, and .elem
def lookup(t,i):
    if t == null: return false
    if t.elem == i: return true
    else if t.elem < i: return lookup(t.left, i)
    else if t.elem > i: return lookup(t.right, i)
```

Challenge: Implement lookup w/ loops

```
# Assume node(elem,left,right) is a constructor
def add(t,i):
    if t == null: new node(i,null,null)
    if t.elem == i: return t
    else if t.elem < i:
        return node(t.elem,add(t.left,i),t.right)
    else if t.elem > i:
        return node(t.elem,t.right,add(t.right,i))
```

```
# Assume node(elem,left,right) is a constructor
def add(t,i):
    if t == null: new node(i,null,null)
    if t.elem == i: return t
    else if t.elem < i:
        return node(t.elem,add(t.left,i),t.right)
    else if t.elem > i:
        return node(t.elem,t.right,add(t.right,i))
```

Challenge: Implement add w/ loops

Observation: BSTs can store **more than just numbers**

Only need **total ordering** (any two can be compared)

- ➡Strings

- ➡Doubles

- ➡Other user defined types

- ➡Some langs allow overloading <

Can also use as basis for other data structures
(e.g., dictionary: nodes key/value pairs)

insert $O(\text{height})$

lookup $O(\text{height})$

$O(\log(\text{size}))$ when *balanced*


insert $O(\text{height})$

lookup $O(\text{height})$


$O(\log(\text{size}))$ when *balanced*

Naive insertion **does not** balance tree :(

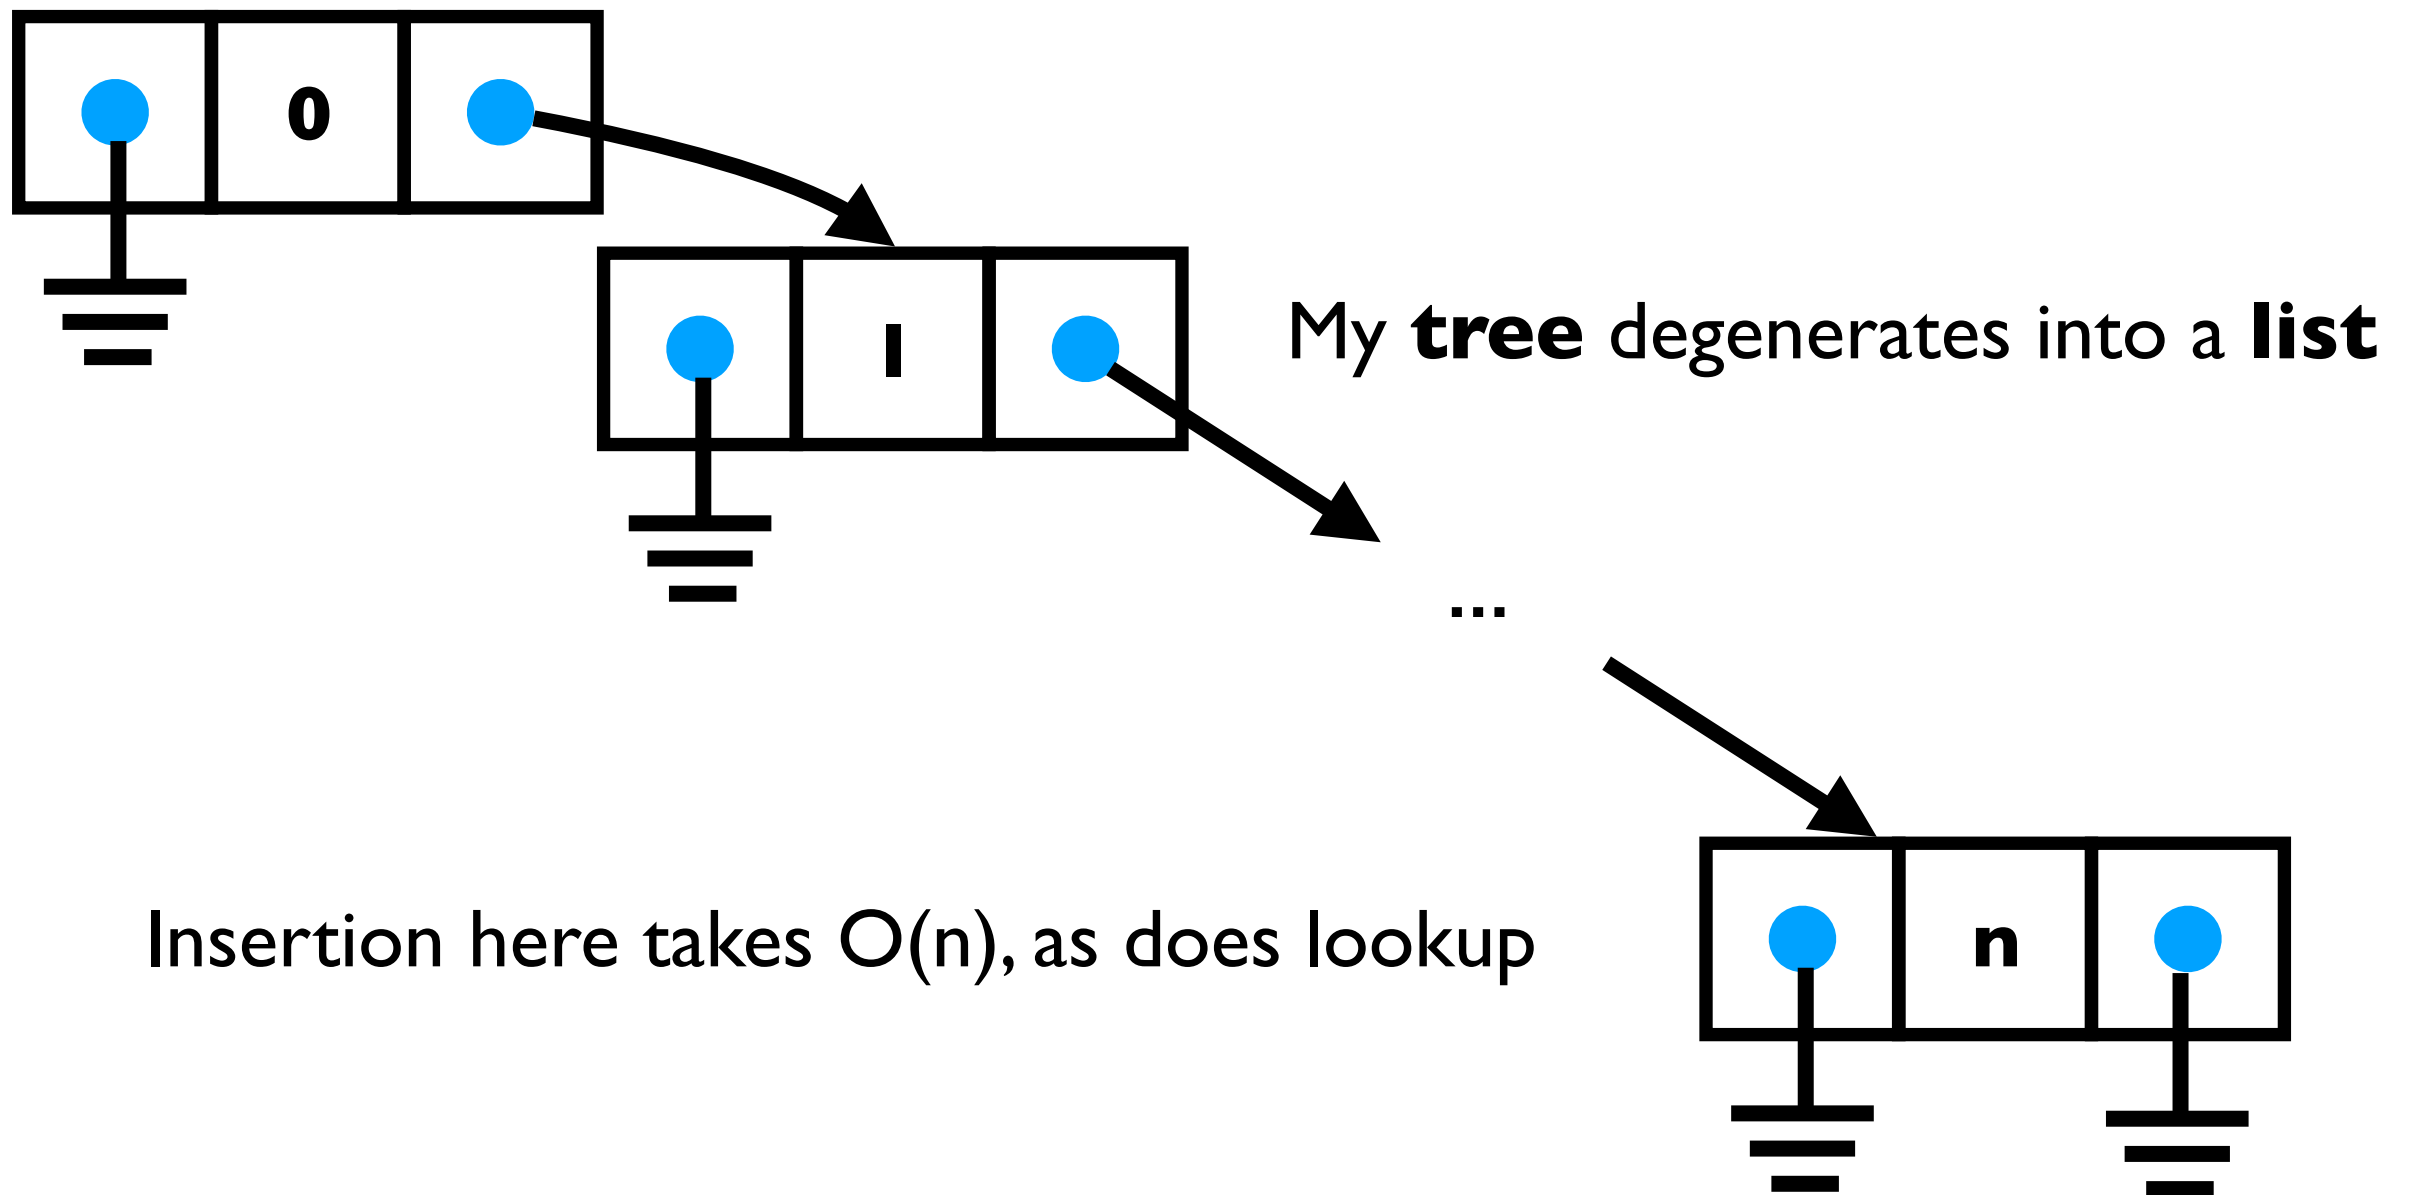
Let's say I start with a 1-element tree...



Then extend it...



Generally: inserting in sorted order is **bad**




Question

Can we ensure good performance generally?


- Precompute **best** BST (dynamic programming)
- **Randomize** insertion order
- Build *even smarter* data structures:
 - Red-Black trees maintain “balanced-ish” trees
 - AVL trees “rebalance” the tree

Balanced Binary Trees



Almost as much stuff on left as right

Balanced Binary Trees




Definition. A tree is “height-balanced” if:

- For each subtree
 - The height of the left subtree is within 1 of the right subtree

Definition. A tree is “height-balanced” if:

- For each subtree
 - The height of the left subtree is within 1 of the right subtree




Which of {A,B,C} are balanced?

Claim (Unproven): If you're using a height-balanced tree, lookups are $O(\log(\text{height}))$


Observation: Inserting into a tree can cause it to become unbalanced

Trick: “Rebalance” the tree upon insertion

Note: I won't ask questions about AVL trees / rebalancing on exam (but I might ask questions about whether trees are height-balanced)




Inserting 10 is ok...




Unbalanced!


Inserting 20 throws off balance of root




Trick: “Rebalance” the tree



Trick: “Rebalance” the tree




Trick: “Rebalance” the tree




Valid b/c this is still a BST!


This is called “rotation”




This is a Right-Right (RR) Rotation




Also need to consider RL rotation



And LL rotation



Last: LR rotation




Generally: AVL trees

- AVL trees are rebalancing binary trees that use rotations to ensure balance invariants
- Generalizes these cases but this is the basic idea
- To insert:
 - Perform BST insertion and then...
 - Go “back up” the spine balancing along the way
- $O(\log(\text{height}))$ performance w/ higher constant factors
 - Rebalancing a node constant time

Other options too..

- Red/Black trees:
 - “Colors” each node either red or black
 - Root is black
 - Every red node’s children must be black
 - Never two black nodes in a row
 - Less balanced, faster insertion, slower lookup



Observation

- Both red-black and AVL trees are **imperative**
- Rebalancing is an inherently imperative operation
 - Changes structure of tree
- Other ultra-fancy data structures fix some of this:
 - E.g., Hash Array-Mapped Trie (HAMT)
 - Will possibly see this later in class...

List

insert

$O(1)$



Simple

lookup

$O(n)$



Insertions frequent



Lookups frequent

List

insert

$O(1)$



lookup

$O(n)$



Simple

Insertions frequent

Lookups frequent

Sorted Array

Also allocates lots of memory

insert

$O(n)$



lookup

$O(\log(n))$



Lookups frequent

Insertions frequent

List

insert

$O(1)$



Simple

lookup

$O(n)$



Insertions frequent



Lookups frequent

Sorted Array

Also allocates lots of memory

insert

$O(n)$



Lookups frequent

lookup

$O(\log(n))$



Insertions frequent

Balanced Binary Tree

balanced

insert

$\sim O(\log(n))$



Lookups frequent



Insertions frequent

lookup

$\sim O(\log(n))$



Maintaining balance hard

Dictionaries

Definition: Dictionary

A dictionary is a key / value mapping

You can think of it as a mathematical function

Key \rightarrow Value

Two main operations

set(Key, Value)

get(Key) \rightarrow Value

set(Key, Value)

get(Key) -> Value

This is the ADT of a dictionary


(Abstract Data Type)

How do we implement it?

(Many possible ways!!)

Implementation I: Association Lists

Key idea: Store a list of **pairs** of keys and values



(In groups...)

How would you implement insert / lookup?

What are their running times?

Are your operations imperative or persistent?

Implementation 2: Lambdas

Key idea: Actually create a function

```
class LambdaDictionary:
    def __init__(self):
        self.f = (lambda key: 1/0)

    def get(self, data):
        return self.f(data)


    def set(self, key, value):
        self.f = (lambda k:
                    value if k == key else self.f(key))
```

Why does this work..? **What is the running time?**

Implementation 3: Balanced BST

Key idea: Each node in BST stores (key,value) pair

Need to order tree in some way (lexicographic order here)



(BTW, lexicographic order essentially means alphabetical order..)

Three Implementations Contrasted

Association List / Functions

Insert	$O(n)$
Lookup	$O(n)$

Balanced BST

Insert	$O(\log(n))$
Lookup	$O(\log(n))$

Where n is number of inserted elements

Next Time: Better Solution via Hash-Tables

Hash tables get us a dictionary with..

Set $\sim O(1)$

Insert $\sim O(1)$

Under appropriate conditions