

And review of last time..

# Return-Oriented Programming



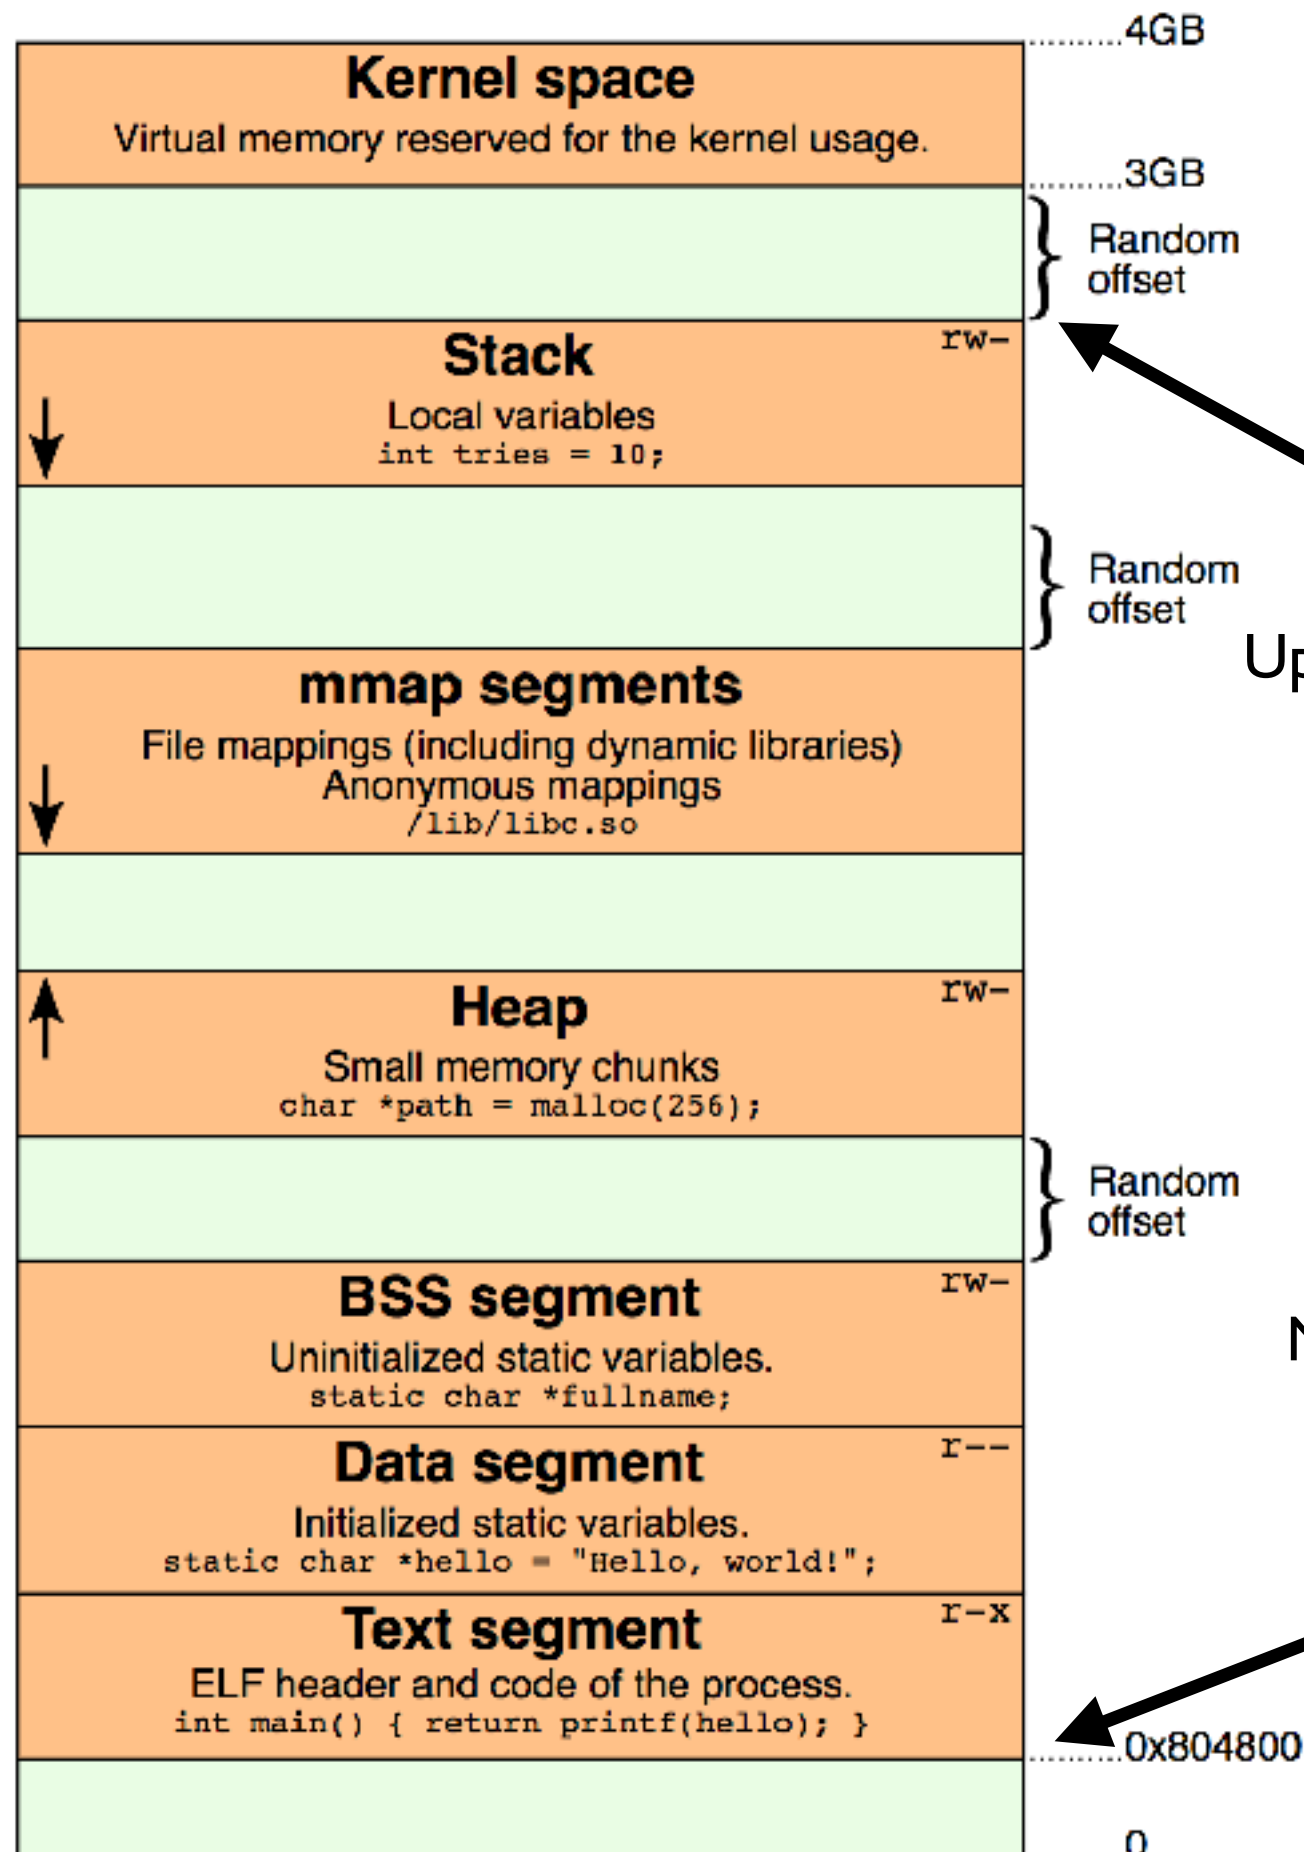
A photograph of two small tabby kittens in a grassy field. One kitten is perched on a tree stump, looking up at another kitten that is jumping or falling through the air above it. The background is a soft-focus green field.

# **Memory Defenses**



**A**dress **S**pace **L**ayout **R**andomization

Randomizes the position of stack, heap, program, libraries



Upshot: Even if you can inject code into the stack, you won't be able to **find** it

Note that the text segment (binary code for program) **isn't** randomized here

# Detour: Position Independent / Relocatable Code

- `.text` segment holds binary representation of program's code
  - All globbed together, each function one after other
- **Within** the text segment, the position of functions **not** changed
  - E.g., if `foo` is at `bar+0x300`, it will **always** be at `bar+0x300`

Program depends on offsets *within* text segment

# Detour: Position Independent / Relocatable Code

- `.text` segment holds binary representation of program's code
  - All globbed together, each function one after other
- **Within** the text segment, the position of functions **not** changed
  - E.g., if `foo` is at `bar+0x300`, it will **always** be at `bar+0x300`

Program depends on offsets *within* text segment

However, **base address** of text could be randomized

- Code must be compiled with a flag `-fPIE`
  - (Position-Independent Execution)

Q: Why **wouldn't** code be compiled with PIE?

A: Can be **faster** to run code that knows its base address

Shows you the **memory maps** for the **current process**

```
cat /proc/self/maps
```

# Exercise

```
micinski@micinski:~$ cat /proc/self/maps
00400000-0040c000 r-xp 00000000 08:01 1704116      /bin/cat
0060b000-0060c000 r--p 0000b000 08:01 1704116      /bin/cat
0060c000-0060d000 rw-p 0000c000 08:01 1704116      /bin/cat
00d37000-00d58000 rw-p 00000000 00:00 0          [heap]
7fb458920000-7fb458bf8000 r--p 00000000 08:01 2635826  /usr/lib/locale/locale-archive
7fb458bf8000-7fb458db8000 r-xp 00000000 08:01 25562894  /lib/x86_64-linux-gnu/libc-2.23.so
7fb458db8000-7fb458fb8000 ---p 001c0000 08:01 25562894  /lib/x86_64-linux-gnu/libc-2.23.so
7fb458fb8000-7fb458fbc000 r--p 001c0000 08:01 25562894  /lib/x86_64-linux-gnu/libc-2.23.so
7fb458fbc000-7fb458fbe000 rw-p 001c4000 08:01 25562894  /lib/x86_64-linux-gnu/libc-2.23.so
7fb458fbe000-7fb458fc2000 rw-p 00000000 00:00 0
7fb458fc2000-7fb458fe8000 r-xp 00000000 08:01 25562855  /lib/x86_64-linux-gnu/ld-2.23.so
7fb45919f000-7fb4591c4000 rw-p 00000000 00:00 0
7fb4591e5000-7fb4591e7000 rw-p 00000000 00:00 0
7fb4591e7000-7fb4591e8000 r--p 00025000 08:01 25562855  /lib/x86_64-linux-gnu/ld-2.23.so
7fb4591e8000-7fb4591e9000 rw-p 00026000 08:01 25562855  /lib/x86_64-linux-gnu/ld-2.23.so
7fb4591e9000-7fb4591ea000 rw-p 00000000 00:00 0
7fff36194000-7fff361b5000 rw-p 00000000 00:00 0          [stack]
7fff361f8000-7fff361fa000 r--p 00000000 00:00 0          [vvar]
7fff361fa000-7fff361fc000 r-xp 00000000 00:00 0          [vdso]
ffffffffffffff600000-ffffffffffffff601000 r-xp 00000000 00:00 0          [vsyscall]
```

Find text, static app data, and app global variables



Text segment (Read+Execute)

Data segment (Read)

Global variables (Read+Write)

micinski@micinski:~\$ cat /proc/self/maps

```
00400000-0040c000 r-xp 00000000 08:01 1704116
0060b000-0060c000 r--p 0000b000 08:01 1704116
0060c000-0060d000 rw-p 0000c000 08:01 1704116
00d37000-00d58000 rw-p 00000000 00:00 0
7fb458920000-7fb458bf8000 r--p 00000000 08:01 2635826
7fb458bf8000-7fb458db8000 r-xp 00000000 08:01 25562894
7fb458db8000-7fb458fb8000 ---p 001c0000 08:01 25562894
7fb458fb8000-7fb458fbc000 r--p 001c0000 08:01 25562894
7fb458fbc000-7fb458fbe000 rw-p 001c4000 08:01 25562894
7fb458fbe000-7fb458fc2000 rw-p 00000000 00:00 0
7fb458fc2000-7fb458fe8000 r-xp 00000000 08:01 25562855
7fb45919f000-7fb4591c4000 rw-p 00000000 00:00 0
7fb4591e5000-7fb4591e7000 rw-p 00000000 00:00 0
7fb4591e7000-7fb4591e8000 r--p 00025000 08:01 25562855
7fb4591e8000-7fb4591e9000 rw-p 00026000 08:01 25562855
7fb4591e9000-7fb4591ea000 rw-p 00000000 00:00 0
7fff36194000-7fff361b5000 rw-p 00000000 00:00 0
7fff361f8000-7fff361fa000 r--p 00000000 00:00 0
7fff361fa000-7fff361fc000 r-xp 00000000 00:00 0
ffffffffffffff600000-ffffffffffffff601000 r-xp 00000000 00:00 0
```

/bin/cat

/bin/cat

/bin/cat

[heap]

/usr/lib/locale/locale-archive

/lib/x86\_64-linux-gnu/libc-2.23.so

/lib/x86\_64-linux-gnu/libc-2.23.so

/lib/x86\_64-linux-gnu/libc-2.23.so

/lib/x86\_64-linux-gnu/libc-2.23.so

/lib/x86\_64-linux-gnu/ld-2.23.so

/lib/x86\_64-linux-gnu/ld-2.23.so

/lib/x86\_64-linux-gnu/ld-2.23.so

[stack]

[vvar]

[vdso]

[vsyscall]

# Defeating ASLR

Two main methods: **brute force** and **derandomization**

Just try a bunch of different addresses and hope for the best

(Doesn't work so well in a 64-bit address space..)

# Defeating ASLR

Two main methods: **brute force** and **derandomization**

Get program to **leak** the value of a pointer to you

# Exercise: break this program

```
void insecure(char *str) {  
    char buffer[100];  
    if (str[3] == 'H') {  
        send("&x", &buffer); // Assume this goes back to user  
    }  
    strcpy(buffer, str);  
}
```

# Exercise: break this program

```
void insecure(char *str) {  
    char buffer[100];  
    if (str[3] == 'H') {  
        send("&x", &buffer); // Assume this goes back to user  
    }  
    strcpy(buffer, str);  
}
```

This example is obviously fake

However, much more common is **error logs**

(If you can convince an app to throw an error to you that contains pointer, you win!)



<https://fail0verflow.com/blog/2017/ps4-crashdump-dump/>

**PS4 Kernel dumped in 11 days via error logs attacker can control!**

# Careful: learning address of stack doesn't tell you where text segment is

```
micinski@micinski:~$ cat /proc/self/maps
```

00400000-0040c000	r-xp	00000000	08:01	1704116	/bin/cat
0060b000-0060c000	r--p	0000b000	08:01	1704116	/bin/cat
0060c000-0060d000	rw-p	0000c000	08:01	1704116	/bin/cat
00d37000-00d58000	rw-p	00000000	00:00	0	[heap]
7fb458920000-7fb458bf8000	r--p	00000000	08:01	2635826	/usr/lib/locale/locale-archive
7fb458bf8000-7fb458db8000	r-xp	00000000	08:01	25562894	/lib/x86_64-linux-gnu/libc-2.23.so
7fb458db8000-7fb458fb8000	---p	001c0000	08:01	25562894	/lib/x86_64-linux-gnu/libc-2.23.so
7fb458fb8000-7fb458fbc000	r--p	001c0000	08:01	25562894	/lib/x86_64-linux-gnu/libc-2.23.so
7fb458fbc000-7fb458fbe000	rw-p	001c4000	08:01	25562894	/lib/x86_64-linux-gnu/libc-2.23.so
7fb458fbe000-7fb458fc2000	rw-p	00000000	00:00	0	
7fb458fc2000-7fb458fe8000	r-xp	00000000	08:01	25562855	/lib/x86_64-linux-gnu/ld-2.23.so
7fb45919f000-7fb4591c4000	rw-p	00000000	00:00	0	
7fb4591e5000-7fb4591e7000	rw-p	00000000	00:00	0	
7fb4591e7000-7fb4591e8000	r--p	00025000	08:01	25562855	/lib/x86_64-linux-gnu/ld-2.23.so
7fb4591e8000-7fb4591e9000	rw-p	00026000	08:01	25562855	/lib/x86_64-linux-gnu/ld-2.23.so
7fb4591e9000-7fb4591ea000	rw-p	00000000	00:00	0	
7fff36194000-7fff361b5000	rw-p	00000000	00:00	0	[stack]
7fff361f8000-7fff361fa000	r--p	00000000	00:00	0	[vvar]
7fff361fa000-7fff361fc000	r-xp	00000000	00:00	0	[vdso]
ffffffffffffff600000-ffffffffffffff601000	r-xp	00000000	00:00	0	[vsyscall]

**N**on **eX**ecutable (stack / heap)

W^X is a simple concept: don't let the programmer execute parts of memory that they can also write

Simple and Effective Defense!

Coordinate w/ CPU

# Defeating NX / W<sup>X</sup>:



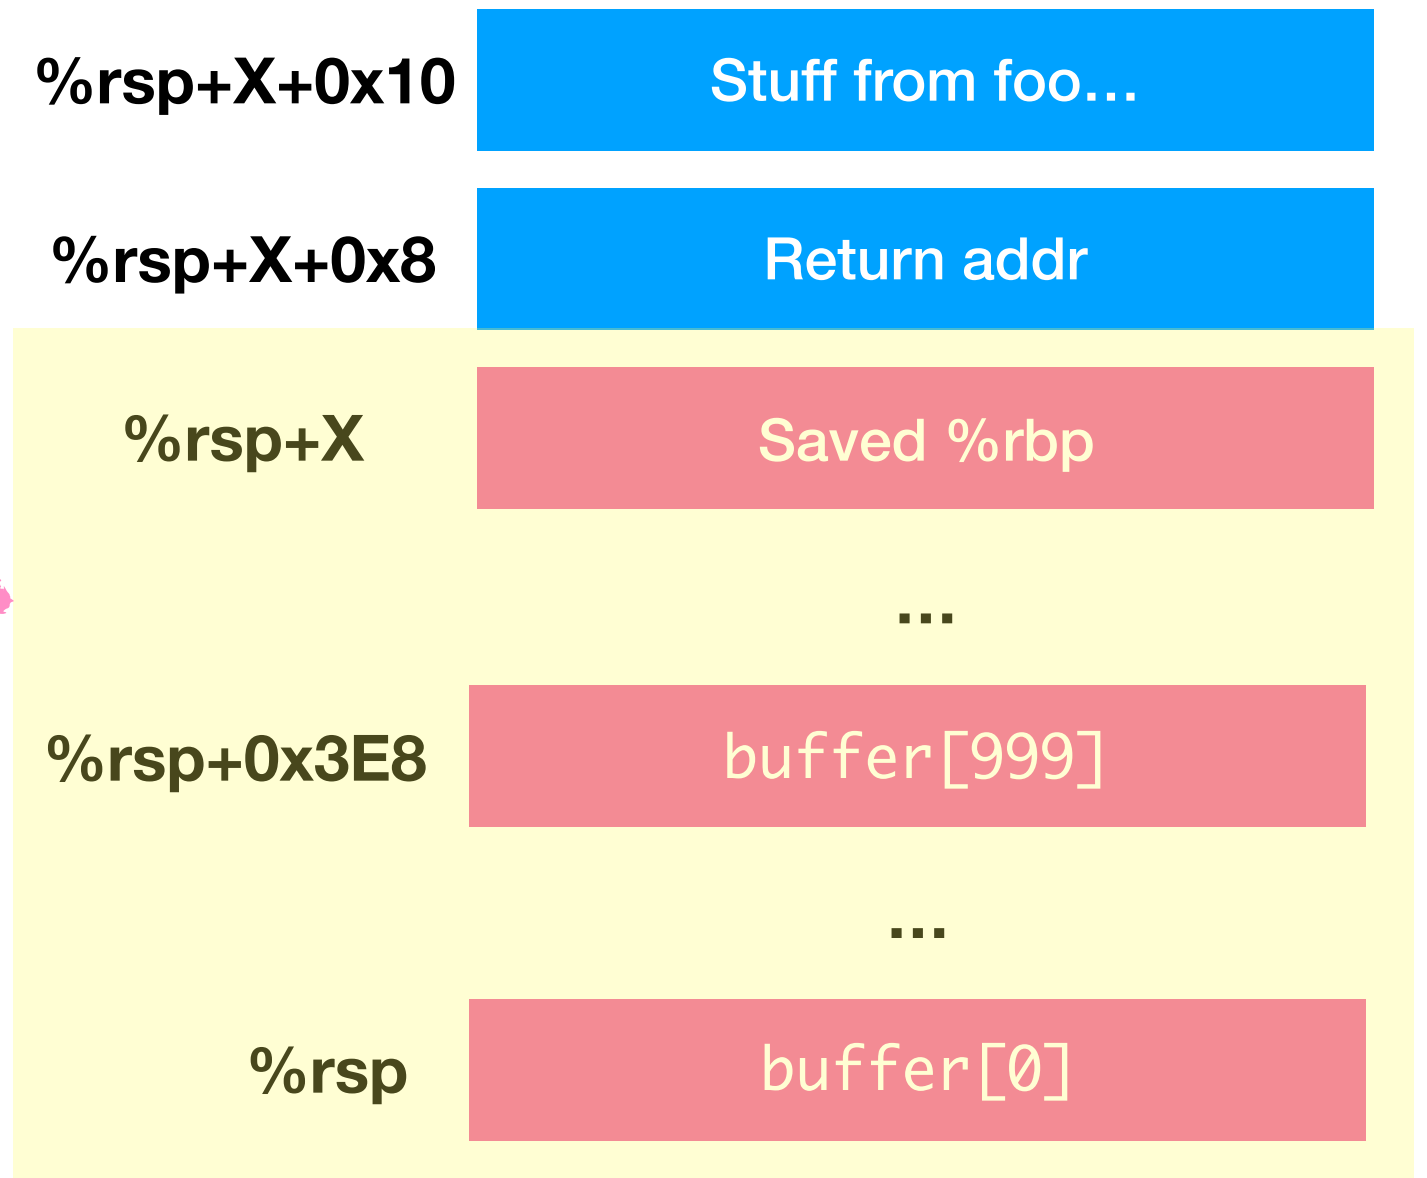

Return-to-libc



Return-oriented-programming

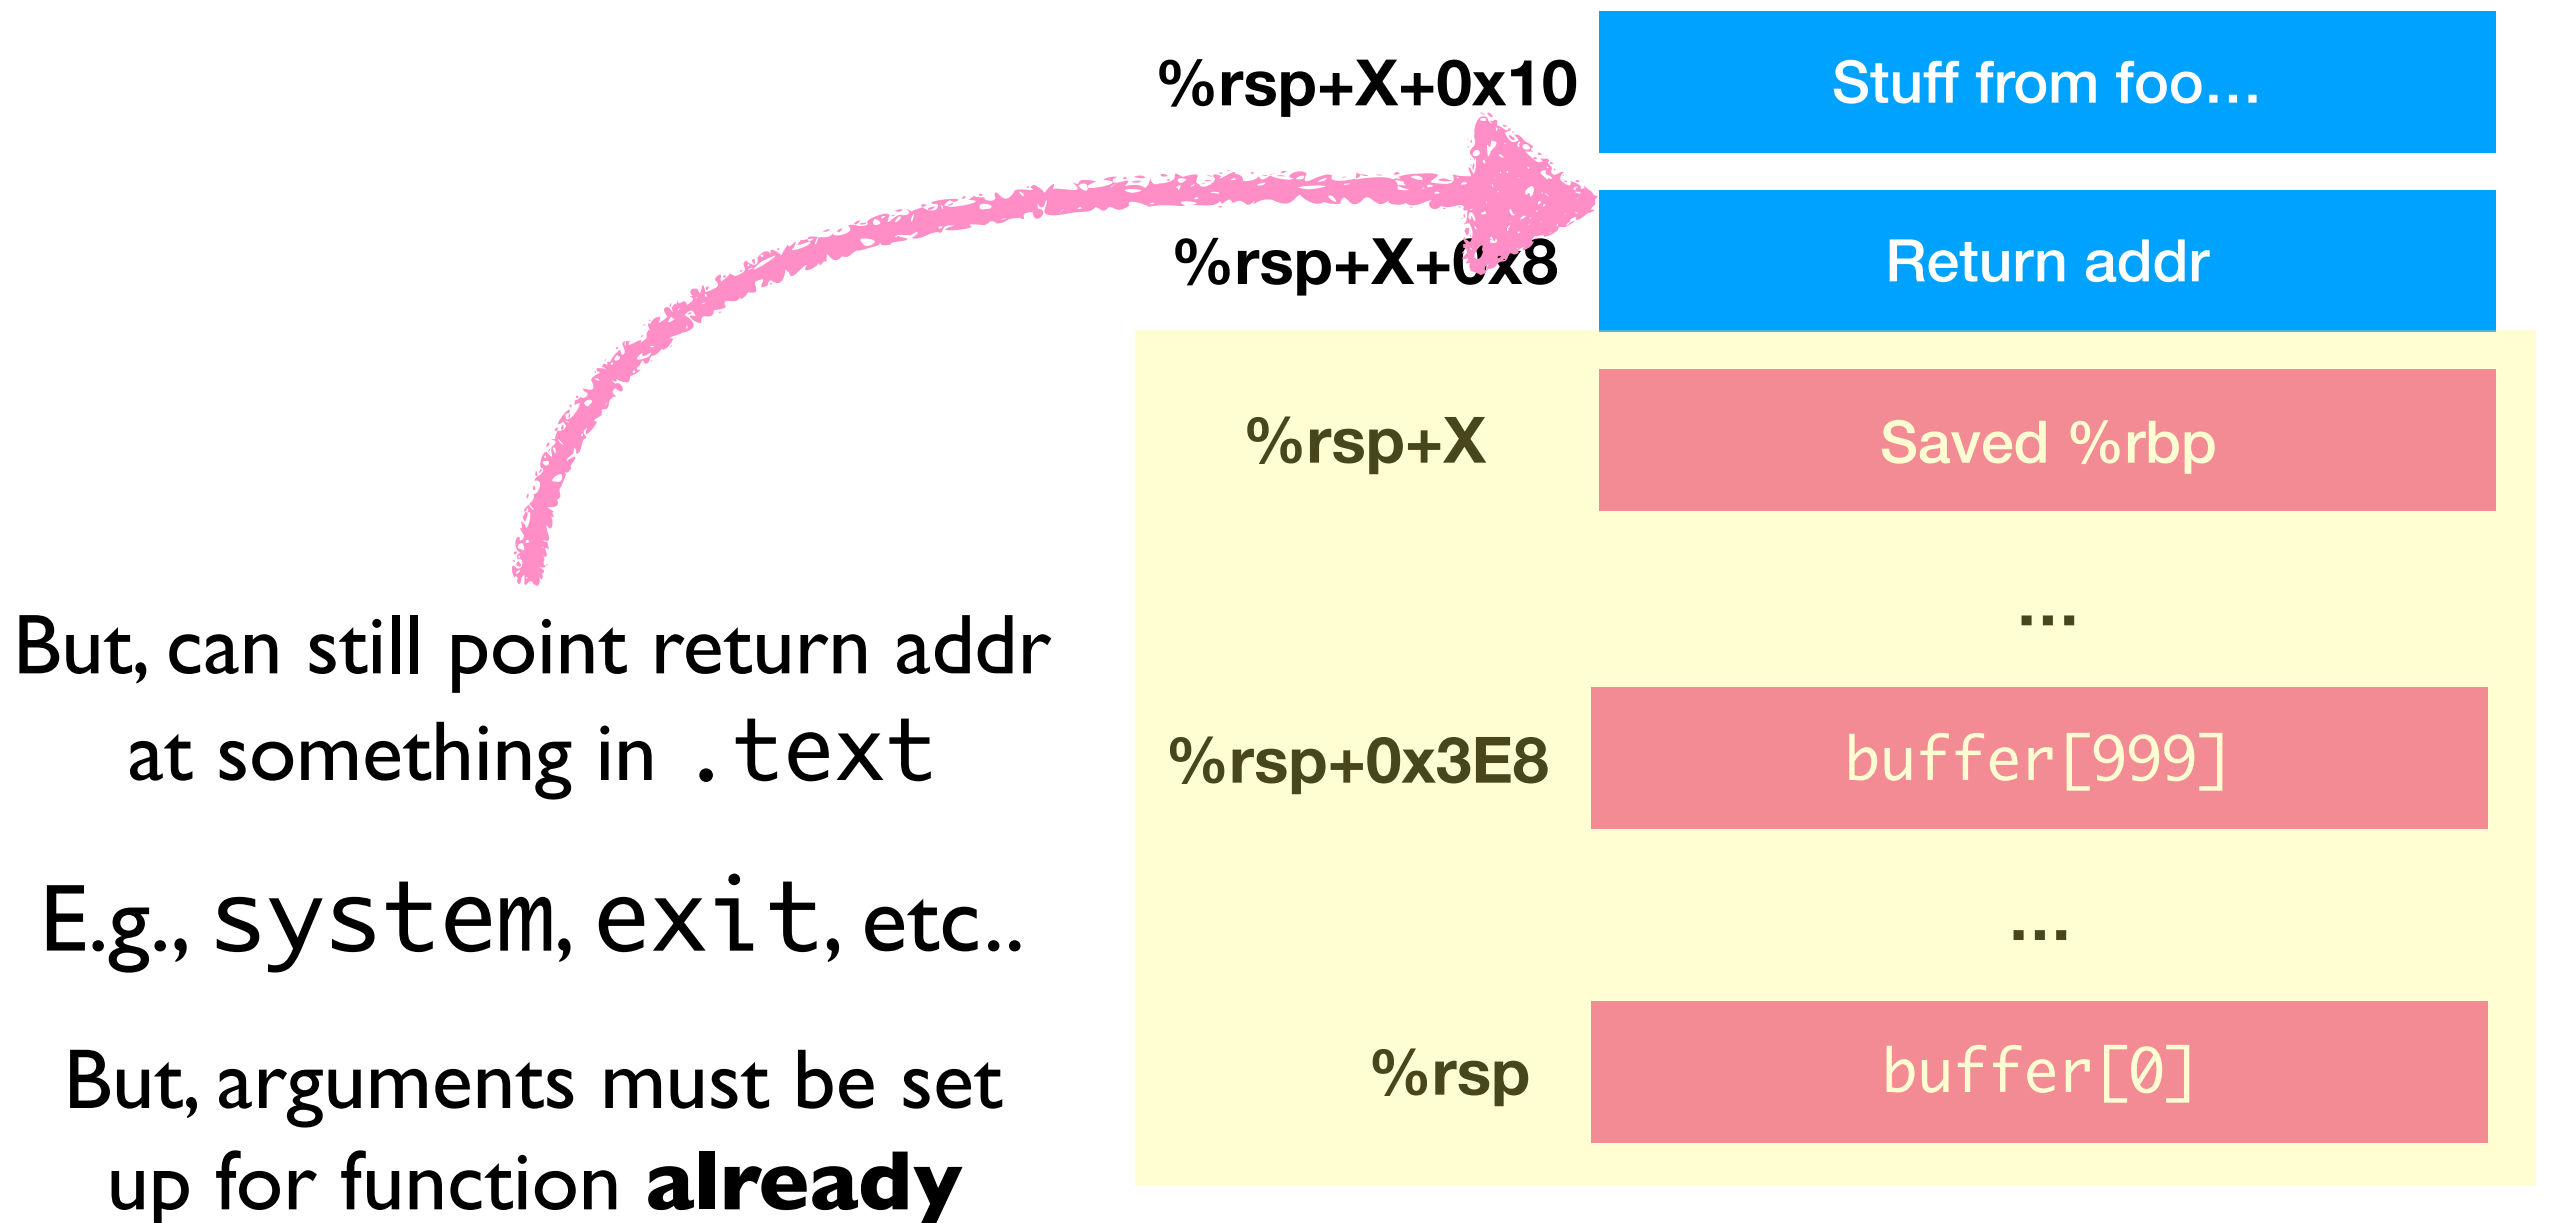
# Return-to-libc

NX: If we try to execute shellcode here, program will **crash!**





# Return-to-libc





# Stack Canaries

Idea: use a **known value** that—if it gets smashed over—alerts you to presence

# “Normal” execution

**%rsp+X+0x10**

Stuff from foo...

**%rsp+X+0x8**

Return addr

**%rsp+X**

Saved %rbp

...

**%rsp+0x3E8**

buffer[999]

...

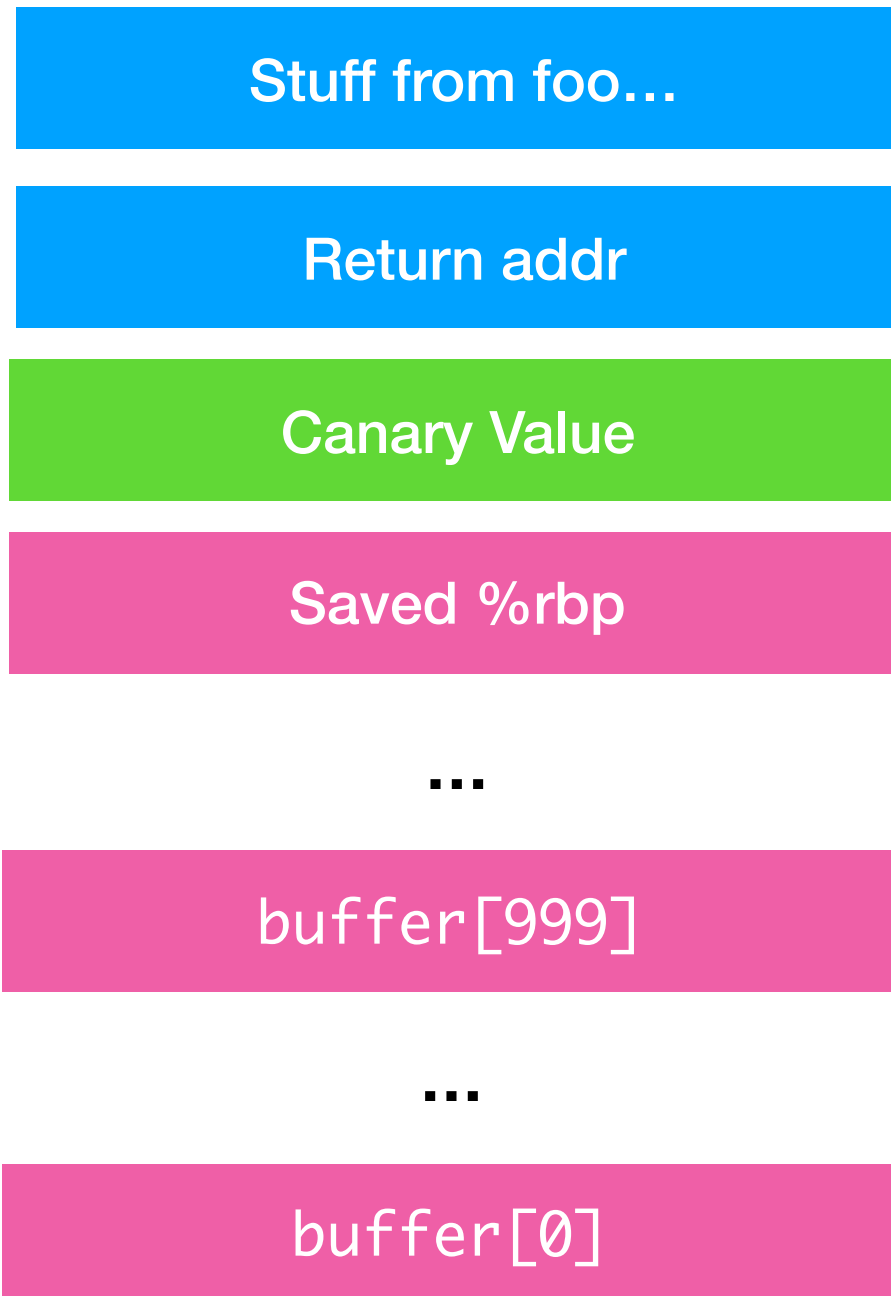
**%rsp**

buffer[0]

# Canary Insertion

**Compiler Inserts  
This Canary**  
(Upon function entry)

Before exiting, **check**  
canary to ensure same



**Exercise:** Compile with and *without* `-fno-stack-protector`



# Defeating Canaries

Can still “skip past” canary occasionally

If attacks “owns” x, can set to skip canary

```
void foo(char *p, int x) {  
    char buffer[100];  
    strcpy(buffer+x,p);  
}
```

# Defeating Canaries

Even if stack overflows can't happen,  
heap overflows can...

```
struct closure {  
    int x;  
    int y;  
    void (*f)(int);  
    char str[8];  
}  
  
closure *x =  
    malloc(sizeof(closure));  
strcpy(x->str, owned_string);  
x->f(42);
```

**Exercise:** Describe w/ partner how you would break **this** program

```
struct closure {  
    int x;  
    int y;  
    char str[100];  
    void (*f)(int);  
}
```

```
int main(int argc, char **argv) {  
    closure *x =  
        malloc(sizeof(closure));  
    strcpy(x->str, argv[1]);  
    x->f(42);  
}
```

In practice, **many** of these defenses  
are employed, and they really do  
**pretty well**

However, the thinking here builds intuition  
for things we still see today...

# Return-Oriented- Programming

Way of “scavenging” through the program’s binary code to trick it into doing **what you want**



Say I wanted to do the following:

- Set %rax to 0
- Execute the “syscall” instruction

@Kris: Write this on board

Say I wanted to do the following:

- Set %rdi to 1 (arg for exit)
- Set %rax to 60 (exit)
- Execute the “syscall” instruction

If I have NX turned on, I can't just **inject** this into the program:

```
movq $1, %rdi
movq $60, %rax
syscall
```

**What might I do instead?**

What might I do instead?


I could try to see the program already has a function  
that does this already and use that.

(I.e., return-to-libc)

What might I do instead?


What if I can't find a whole function that does this?

Normally... Function starts here  
and continues until (either) **ret**




```
0xF000: pushq %rbp
0xF002: movq %rsp, %rbp
0xF004: subq $12, %rsp
0xF007: mov %eax, -4(%rbp)
0xF009: mov %eax, -8(%rbp)
0xF00b: mov %eax, -12(%rbp)
0xF00e: add %eax, %eax
0xF010: compl %eax, %eax
0xF013: jmpg 0xF01d
0xF015: addq $12, %rsp
0xF018: leave
0xF019: mov $60, %rax
0xF01b: syscall
0xF01c: ret
0xF01d: addq $12, %rsp
0xF01f: leave
0xF020: ret
```

Normally... Function starts here  
and continues until (either) **ret**



```
0xF000: pushq %rbp
0xF002: movq %rsp, %rbp
0xF004: subq $12, %rsp
0xF007: mov %eax, -4(%rbp)
0xF009: mov %eax, -8(%rbp)
0xF00b: mov %eax, -12(%rbp)
0xF00e: add %eax, %eax
0xF010: compl %eax, %eax
0xF013: jmpg 0xF01d
0xF015: addq $12, %rsp
0xF018: leave
0xF019: mov $60, %rax
0xF01b: syscall
0xF01c: ret
0xF01d: addq $12, %rsp
0xF01f: leave
0xF020: ret
```

But nothing stops me from  
jumping right **here!**



So I could look through binary and find all places with **ret** and jump to any number of bytes before that.

Observation: can execute sequences of code that weren't **technically** in program to begin with



Observation: x86\_64 instructions are **variable length**

Like words...

“the address”

Observation: x86\_64 instructions are **variable length**

Like words...

“the **address**”

Observation: x86\_64 instructions are **variable length**


Like words...

**“the address”**

```
f7 c7 07 00 00 00
0f 95 45 c3
```

```
test $0x00000007 %edi
setnzb -61(%ebp)
```

Read starting at c7



```
c7 07 00 00 00 0f
95
45
c3
```

```
movl $0xf000000, (%edi)
xchg %ebp, %eax
inc %ebp
ret
```

Let's say that I want to call D01F and **then** F019

```
...  
0xD01F: pop %rdi  
0xD020: ret  
...
```

```
...  
0xF019: mov $60, %rax  
0xF01B: syscall  
0xF01C: ret  
...
```

To “set up” the attack we put 0xD01F in saved RIP

...

0xD01F: pop %rdi

0xD020: ret

...

...


0xF019: mov \$60, %rax

0xF01B: syscall

0xF01C: ret

...

Bar's  
frame



Stuff from foo...

Return addr

Saved %rbp


...

buffer[999]

...

buffer[0]

To “set up” the attack we put 0xD01F in saved RIP



...

0xD01F: pop %rdi

0xD020: ret

...

...


0xF019: mov \$60, %rax

0xF01B: syscall

0xF01C: ret

...

Bar's  
frame



Stuff from foo...

0xD01F

Saved %rbp

...

buffer[999]

...

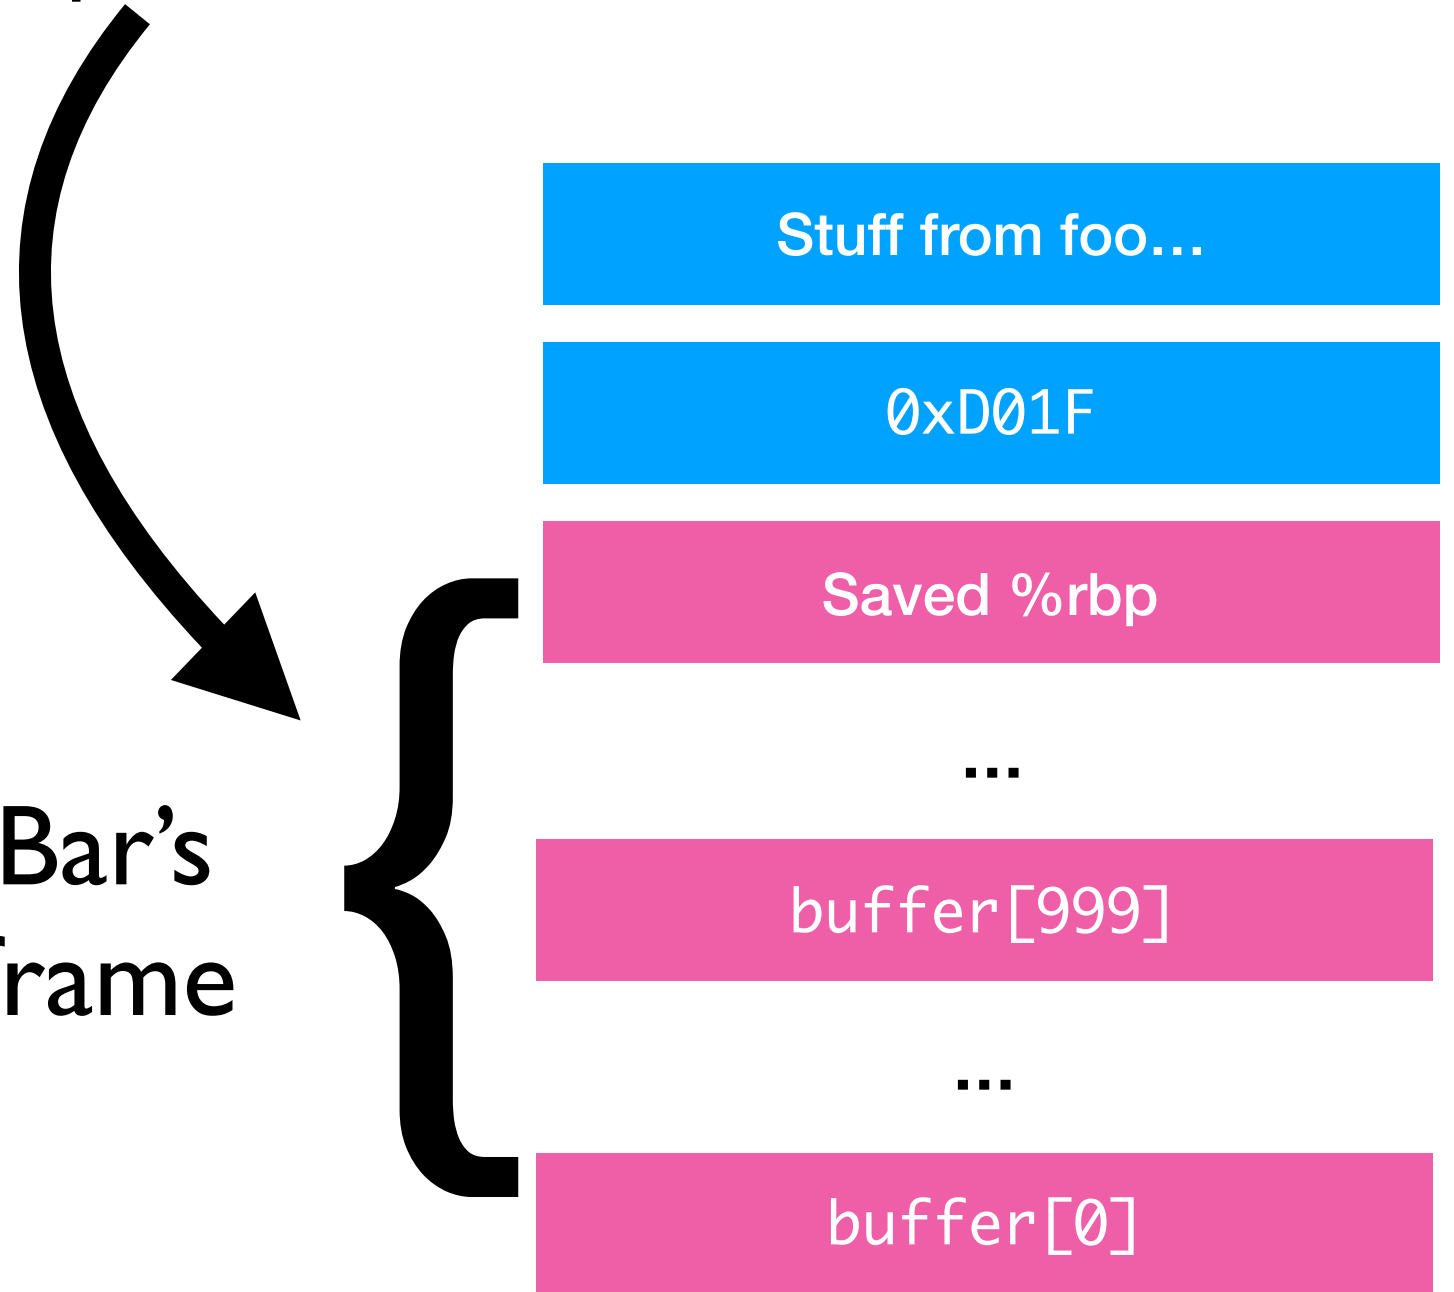
buffer[0]

Before **foo** returns, it pops all of this stuff from the stack

...  
0xD01F: pop %rdi  
0xD020: ret  
...

...  
0xF019: mov \$60, %rax  
0xF01B: syscall  
0xF01C: ret  
...

Bar's  
frame





Now it goes here

...  
0xD01F: pop %rdi  
0xD020: ret  
...

Stuff from foo...

0xD01F

...  
0xF019: mov \$60, %rax  
0xF01B: syscall  
0xF01C: ret  
...

(Rather than it's caller **foo**)

**Super Critical: pops 0xD01F from stack!**

%rsp

Stuff from foo...

...

0xD01F: pop %rdi

0xD020: ret

...

...

0xF019: mov \$60, %rax


0xF01B: syscall

0xF01C: ret

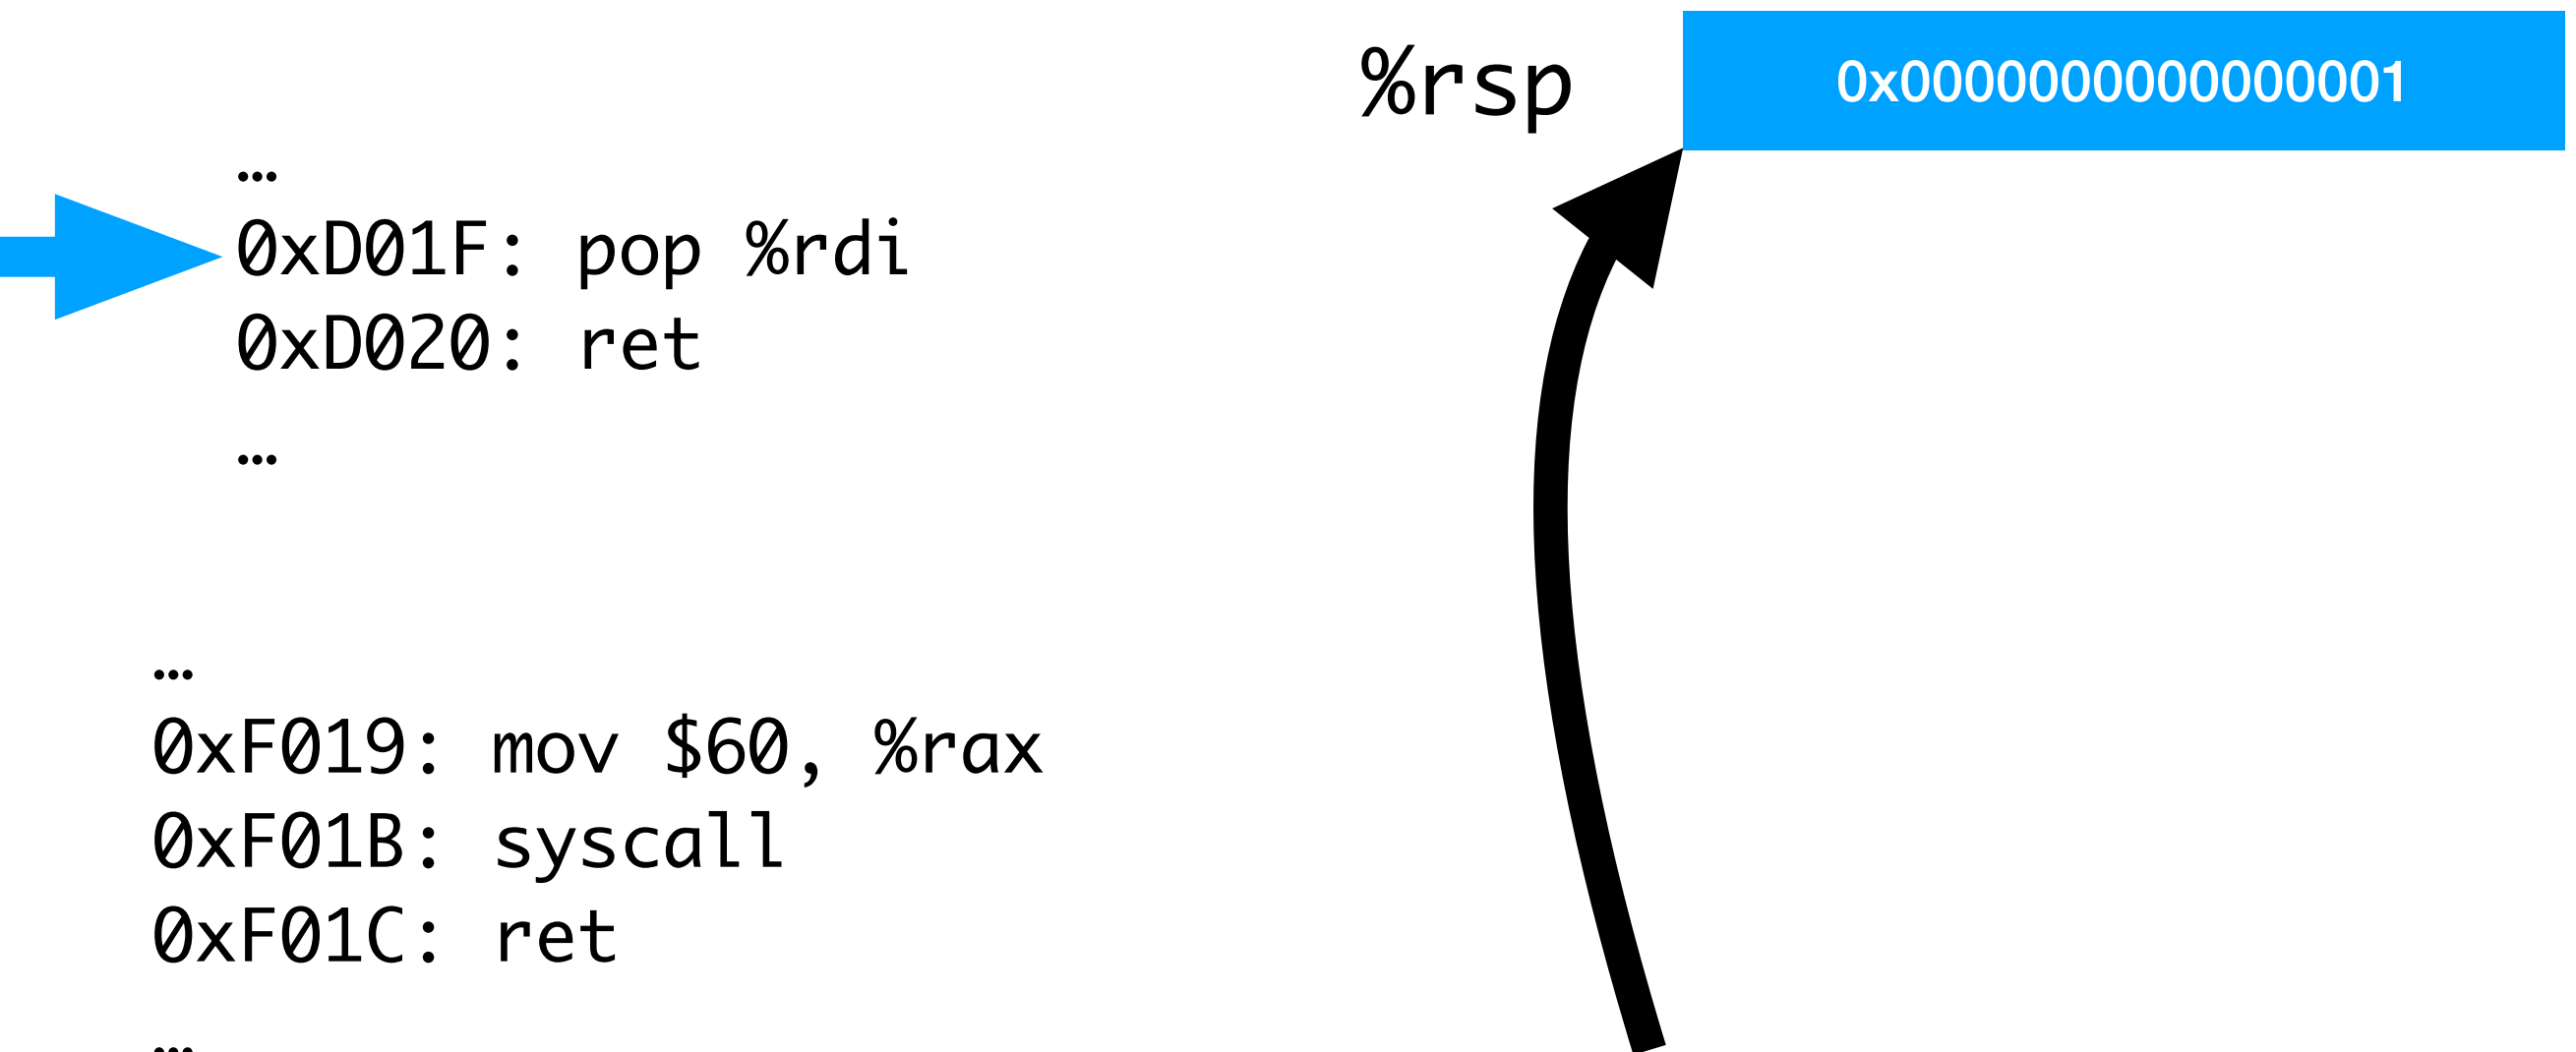
...

So **now** whatever's on stack will be  
popped into %rdi

(Which is previously stuff in **foo**'s stack)



So if I want to put I in RDI, I put it **here**  
(Which is previously stuff in **foo**'s frame)



So if I want to put I in RDI, I put it **here**  
(Which is previously stuff in **foo**'s frame)

%rsp

....

...

0xD01F: pop %rdi

0xD020: ret

...

Now, when the code hits **this** point,  
it's going to execute a return

...

0xF019: mov \$60, %rax

0xF01B: syscall

0xF01C: ret

...

Which will **yet again** go to  
whatever address is in %rsp

%rsp

0xF019

...

0xD01F: pop %rdi

0xD020: ret

...

...

0xF019: mov \$60, %rax

0xF01B: syscall

0xF01C: ret

...

Critical observation: if %rsp is **now**  
0xF019, we'll get what we want

%rsp

0xF019

...

0xD01F: pop %rdi

0xD020: ret

...

...

0xF019: mov \$60, %rax

0xF01B: syscall

0xF01C: ret

...

Critical observation: if %rsp is **now**  
0xF019, we'll get what we want

- Set %rdi to 1 (arg for exit)
- Set %rax to 60 (exit)
- Execute the “syscall” instruction

...

0xD01F: pop %rdi

0xD020: ret

...

...

0xF019: mov \$60, %rax

0xF01B: syscall

0xF01C: ret

...

Critical observation: if %rsp is **now**  
0xF019, we'll get what we want



**Observation:** We can **chain** multiple sequences  
(that all end in **ret**) by setting up the stack right

# Exercise

```
write(1, "Hello, world!", 13);
```

```
%rax = 1 %rdi = 1 %rsi = &"Hello, world", %rdx = 13
```

```
0xC110: pop %rsi  
0xC112: ret
```

```
0x1029: pop %edx  
0x102a: ret
```

```
0xD235: xchang %rdx, %rdi  
0xD238: ret
```

```
0xF019: pop %eax  
0xF01B: ret
```

```
0xB0FF: pop %rdx  
0xB102: ret
```

```
0xCA2F: syscall
```

**Assume  
this is 128**

**buffer = 0x40000**

Saved %rbp

...

buffer[999]

...

buffer[0]